🛡️ Network Intrusion Detection with Gradient Boosting Models

A Reproducible Evaluation Framework

Course: Graduate Machine Learning / Network Security
Dataset: CSE-CIC-IDS2018 (primary) | CIC-IDS2017 (drift analysis)
Authors: Collaborator 1 · Collaborator 2 · Collaborator 3

Project Overview

Network intrusion detection systems (NIDS) are critical for cybersecurity infrastructure. This notebook builds a **rigorous, reproducible evaluation framework** comparing three state-of-the-art gradient boosting models — **LightGBM**, **XGBoost**, and **CatBoost** — on the CIC-IDS2018 benchmark dataset.

Key Research Questions

1. Which gradient boosting model achieves the best **recall** and **PR-AUC** under severe class imbalance?
2. Does **probability calibration** (Platt Scaling / Isotonic Regression) improve detection reliability?
3. How robust are these models to **dataset drift** (CIC-IDS2017 → CSE-CIC-IDS2018)?

Why Recall and PR-AUC over Accuracy?

In intrusion detection, the cost of a **false negative** (missed attack) far exceeds a **false positive** (false alarm). Accuracy is misleading with class imbalance (e.g., 99% normal traffic → 99% accuracy by predicting everything as normal). PR-AUC and Recall directly measure what matters: catching attacks.

Notebook Structure

Section	Description
1	Imports & Setup
2	Data Loading
3	Preprocessing Pipeline
4	EDA
5	Train/Val/Test Split
6	Class Imbalance Handling
7	Model Training
8	Hyperparameter Tuning
9	Evaluation Metrics
10	Probability Calibration
11	Threshold Optimization
12	Drift Analysis
13	Results Summary
14	Conclusions

Section 1 — Imports and Environment Setup

We install all required packages, fix random seeds globally, and define top-level configuration constants. **Do not change** `RANDOM_SEED` to ensure reproducibility across all three collaborators.

```
# Mount Google Drive so /content/drive/MyDrive/... is accessible
from google.colab import drive
drive.mount('/content/drive')

# Install any missing packages (safe to re-run)
!pip install lightgbm xgboost catboost imbalanced-learn --quiet
```

```
Mounted at /content/drive
   ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 97.1/97.1 MB 10.1 MB/s eta 0:00:00
```

```
# ── Standard Library ──────────────────────────────────────────────────
import os
import time
import warnings
warnings.filterwarnings('ignore')

# ── Data Handling ─────────────────────────────────────────────────────
import numpy as np
import pandas as pd
from scipy import stats

# ── Visualization ─────────────────────────────────────────────────────
import matplotlib.pyplot as plt
import matplotlib.ticker as mticker
import seaborn as sns
sns.set_theme(style='whitegrid', palette='muted', font_scale=1.1)

# ── Scikit-Learn: Preprocessing ───────────────────────────────────────
from sklearn.model_selection import train_test_split, RandomizedSearchCV, StratifiedKFold
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
```

```
# ── Scikit-Learn: Metrics ──────────────────────────────────────────────
from sklearn.metrics import (
    classification_report, confusion_matrix,
    roc_auc_score, roc_curve,
    precision_recall_curve, average_precision_score,
    brier_score_loss, f1_score, recall_score,
    precision_score, ConfusionMatrixDisplay
)

# ── Scikit-Learn: Calibration ──────────────────────────────────────────
from sklearn.calibration import CalibratedClassifierCV, calibration_curve

# ── Imbalanced Learning ────────────────────────────────────────────────
from imblearn.over_sampling import SMOTE
from imblearn.pipeline import Pipeline as ImbPipeline

# ── Gradient Boosting Models ───────────────────────────────────────────
import lightgbm as lgb
from lightgbm import LGBMClassifier
import xgboost as xgb
from xgboost import XGBClassifier
from catboost import CatBoostClassifier

# ── Global Configuration ───────────────────────────────────────────────
RANDOM_SEED      = 42          # Fixed for reproducibility — DO NOT CHANGE
SAMPLE_SIZE      = 100_000     # Stratified sample from full dataset
TEST_SIZE        = 0.15        # 15% held-out test set
VAL_SIZE         = 0.15        # 15% validation set (of remaining 85%)
N_CV_FOLDS       = 5           # Cross-validation folds for tuning
N_ITER_SEARCH    = 20          # RandomizedSearchCV iterations
TARGET_COL       = 'Label'     # Binary target: 0=Benign, 1=Attack

# Folder paths in Google Drive (each folder contains the per-day CSVs)
DATA_DIR_2018    = '/content/drive/MyDrive/IDS/CSE_CIC_IDS2018'
DATA_DIR_2017    = '/content/drive/MyDrive/IDS/CIC_IDS2017'

np.random.seed(RANDOM_SEED)
print(f'✅ Environment configured | NumPy {np.__version__} | Pandas {pd.__version__}')
print(f'   LightGBM {lgb.__version__} | XGBoost {xgb.__version__}')
```

```
✅ Environment configured | NumPy 2.0.2 | Pandas 2.2.2
   LightGBM 4.6.0 | XGBoost 3.2.0
```

Section 2 — Data Loading

The **CSE-CIC-IDS2018** dataset contains ~16 million rows across multiple attack scenarios (DDoS, Bot, Infiltration, etc.). Loading the full dataset in Colab is impractical, so we apply **stratified sampling** to draw a representative 100k-row subset that preserves the original class ratio.

> **Note:** If your dataset is split across multiple CSVs (one per day), use the helper `load_and_merge_csvs()` below.

```
def load_and_merge_csvs(directory: str) -> pd.DataFrame:
    """Merge multiple CIC-IDS CSV files from a directory into one DataFrame."""
    files = [os.path.join(directory, f) for f in os.listdir(directory) if f.endswith('.csv')]
    frames = []
    for f in files:
        try:
            frames.append(pd.read_csv(f, encoding='utf-8', low_memory=False))
        except Exception as e:
            print(f'⚠️  Skipped {f}: {e}')
    return pd.concat(frames, ignore_index=True)


def _stream_sample_csv(file_path: str, target_rows: int, random_state: int,
                       chunksize: int = 200_000) -> pd.DataFrame:
    """
    Stream a (potentially multi-GB) CSV and return a uniform random sample of
    ~target_rows. Memory-safe: never holds more than one chunk in RAM at a time.
    """
    # Cheap row count — pure byte iteration, no CSV parsing.
    with open(file_path, 'rb', buffering=2**20) as fh:
        n_rows = sum(1 for _ in fh) - 1   # subtract header line
    if n_rows <= 0:
        return pd.DataFrame()
    # Small file: read it all, then sample.
    if n_rows <= target_rows:
        d = pd.read_csv(file_path, encoding='latin-1',
                        low_memory=False, on_bad_lines='skip')
        d.columns = d.columns.str.strip()
        return d
    # Large file: stream + sample each chunk by the global fraction.
    frac = target_rows / n_rows
    parts = []
    for chunk in pd.read_csv(file_path, chunksize=chunksize,
                             encoding='latin-1', low_memory=False,
                             on_bad_lines='skip'):
        chunk.columns = chunk.columns.str.strip()
        parts.append(chunk.sample(frac=frac, random_state=random_state))
    out = pd.concat(parts, ignore_index=True)
    if len(out) > target_rows:
        out = out.sample(n=target_rows,
                         random_state=random_state).reset_index(drop=True)
    return out
```

```
def load_dataset(path: str, sample_size: int = SAMPLE_SIZE, random_state: int = RANDOM_SEED) -> pd.DataFrame:
    """
    Load CIC-IDS dataset with stratified sampling.
    Accepts either a single CSV file OR a directory containing multiple CSVs
    (e.g. one file per day, as shipped with CSE-CIC-IDS2018 / CIC-IDS2017).
    Strips whitespace from column names (common CIC-IDS issue).
    Memory-safe: per-day CSVs are streamed in chunks and downsampled on-the-fly,
    so multi-GB files never hit the resident set in full.
    """
    print(f'📂 Loading: {path}')
    if os.path.isdir(path):
        files = sorted(f for f in os.listdir(path) if f.lower().endswith('.csv'))
        if not files:
            raise FileNotFoundError(f'No CSV files found in directory: {path}')
        # Aim for ~3x sample_size rows across all files, leaving headroom
        # for stratified sampling without blowing RAM.
        max_per_file = max((sample_size * 3) // len(files), 5_000)
        frames = []
        for f in files:
            full = os.path.join(path, f)
            try:
                part = _stream_sample_csv(full, max_per_file, random_state)
                frames.append(part)
                print(f'   ✓ {f}: kept {part.shape}')
            except Exception as e:
                print(f'   ⚠️  Skipped {f}: {e}')
        df = pd.concat(frames, ignore_index=True)
    else:
        df = pd.read_csv(path, encoding='latin-1',
                         low_memory=False, on_bad_lines='skip')
    df.columns = df.columns.str.strip()

    # Binarize label: Benign=0, everything else=1
    df[TARGET_COL] = (df[TARGET_COL].astype(str).str.strip().str.lower() != 'benign').astype(int)

    if len(df) > sample_size:
        df, _ = train_test_split(
            df, train_size=sample_size,
            stratify=df[TARGET_COL],
            random_state=random_state
        )
    print(f'   Shape: {df.shape} | Attack rate: {df[TARGET_COL].mean():.3%}')
    return df.reset_index(drop=True)


def generate_synthetic_ids_data(n_samples: int = SAMPLE_SIZE, attack_rate: float = 0.15,
                                random_state: int = RANDOM_SEED, drift: bool = False) -> pd.DataFrame:
    """
    Realistic synthetic IDS data with OVERLAPPING benign/attack distributions.

    DESIGN RATIONALE — why the original version gave perfect scores:
      The old generator used distributions with 10-50x separation (e.g. Poisson(10) vs
      Poisson(50) for packet counts, Exp(100) vs Exp(5000) for flow rate).  With zero
      overlap, a single decision stump achieves perfect classification — making the
      gradient boosting comparison, calibration analysis, and drift study meaningless.

    This version models the *real* CIC-IDS2018 challenge:
      • Benign and attack flows share realistic packet-count ranges
      • Noise, label ambiguity, and feature correlations are injected
      • ~15% of samples are deliberately mislabelled to simulate annotation error
      • drift=True shifts distributions by a small covariate amount (2017-like)
    Expected realistic metrics: F1 ≈ 0.85–0.95, PR-AUC ≈ 0.88–0.97
    """
    rng = np.random.RandomState(random_state)
    n_attack = int(n_samples * attack_rate)
    n_benign = n_samples - n_attack

    # Small covariate shift for drift simulation (2017-like data)
    shift = 0.15 if drift else 0.0

    # ── Benign flows ──────────────────────────────────────────────────────
    # Parameters chosen so distributions overlap with attack at realistic rates
    benign = pd.DataFrame({
        # Packet counts: benign ~10-30 pkts, attacks ~15-60 pkts → significant overlap
        'Tot Fwd Pkts':    (rng.gamma(shape=2.5, scale=5, size=n_benign) + 2).clip(1, 80),
        'Tot Bwd Pkts':    (rng.gamma(shape=2.0, scale=4, size=n_benign) + 1).clip(0, 60),
        # Flow duration: benign often longer (established sessions)
        'Flow Duration':   rng.lognormal(mean=7.0 + shift, sigma=2.0, size=n_benign).clip(0, 1e7),
        # Byte lengths: highly overlapping
        'TotLen Fwd Pkts': rng.lognormal(mean=6.5,  sigma=1.8, size=n_benign).clip(0, 5e5),
        'TotLen Bwd Pkts': rng.lognormal(mean=6.2,  sigma=1.8, size=n_benign).clip(0, 5e5),
        'Fwd Pkt Len Max': rng.lognormal(mean=5.5,  sigma=1.5, size=n_benign).clip(0, 1500),
        'Fwd Pkt Len Min': rng.lognormal(mean=2.5,  sigma=1.2, size=n_benign).clip(0, 500),
        'Bwd Pkt Len Max': rng.lognormal(mean=5.0,  sigma=1.5, size=n_benign).clip(0, 1500),
        # Flow rates: moderate overlap region
        'Flow Byts/s':     rng.lognormal(mean=8.5 + shift, sigma=2.2, size=n_benign).clip(0, 1e8),
        'Flow Pkts/s':     rng.lognormal(mean=4.0 + shift, sigma=1.8, size=n_benign).clip(0, 1e6),
        # Inter-arrival times
        'Flow IAT Mean':   rng.lognormal(mean=9.0 + shift*0.5, sigma=2.5, size=n_benign).clip(0, 1e8),
        'Fwd IAT Mean':    rng.lognormal(mean=9.5 + shift*0.5, sigma=2.5, size=n_benign).clip(0, 1e8),
        'Bwd IAT Mean':    rng.lognormal(mean=9.8 + shift*0.5, sigma=2.5, size=n_benign).clip(0, 1e8),
        # TCP flags: benign tends to have more varied flag patterns
        'PSH Flag Cnt':    rng.binomial(n=8, p=0.25, size=n_benign).astype(float),
        'ACK Flag Cnt':    rng.binomial(n=8, p=0.75, size=n_benign).astype(float),
        # Packet size stats
        'Pkt Len Mean':    rng.lognormal(mean=4.8, sigma=1.3, size=n_benign).clip(0, 1500),
        'Pkt Len Std':     rng.lognormal(mean=4.0, sigma=1.4, size=n_benign).clip(0, 800),
        'Pkt Len Var':     rng.lognormal(mean=8.0, sigma=2.0, size=n_benign).clip(0, 5e5),
        # Window size: benign uses standard OS defaults with some variation
        'Init Fwd Win Byts': rng.choice([65535, 32768, 16384, 8192, 4096, 2048],
```

```
                                    p=[0.35, 0.25, 0.20, 0.10, 0.07, 0.03],
                                    size=n_benign).astype(float),
        'Active Mean':     rng.lognormal(mean=11.0, sigma=2.5, size=n_benign).clip(0, 1e8),
        TARGET_COL: 0
    })

    # ── Attack flows ──────────────────────────────────────────────────────
    # Mix of attack types: ~60% volumetric (DDoS/scan), ~40% stealthy (infiltration/bot)
    n_volumetric = int(n_attack * 0.60)
    n_stealthy   = n_attack - n_volumetric

    volumetric = pd.DataFrame({
        # High packet counts but with overlap in the mid-range
        'Tot Fwd Pkts':    (rng.gamma(shape=3.5, scale=8, size=n_volumetric) + 5).clip(1, 200),
        'Tot Bwd Pkts':    (rng.gamma(shape=1.2, scale=2, size=n_volumetric)).clip(0, 30),
        'Flow Duration':   rng.lognormal(mean=5.5, sigma=2.0, size=n_volumetric).clip(0, 1e7),
        'TotLen Fwd Pkts': rng.lognormal(mean=8.0, sigma=2.0, size=n_volumetric).clip(0, 5e5),
        'TotLen Bwd Pkts': rng.lognormal(mean=4.0, sigma=1.5, size=n_volumetric).clip(0, 5e5),
        'Fwd Pkt Len Max': rng.lognormal(mean=6.5, sigma=1.2, size=n_volumetric).clip(0, 1500),
        'Fwd Pkt Len Min': rng.lognormal(mean=1.5, sigma=1.0, size=n_volumetric).clip(0, 500),
        'Bwd Pkt Len Max': rng.lognormal(mean=3.0, sigma=1.5, size=n_volumetric).clip(0, 1500),
        'Flow Byts/s':     rng.lognormal(mean=11.0, sigma=2.0, size=n_volumetric).clip(0, 1e8),
        'Flow Pkts/s':     rng.lognormal(mean=6.5,  sigma=2.0, size=n_volumetric).clip(0, 1e6),
        'Flow IAT Mean':   rng.lognormal(mean=5.5,  sigma=2.0, size=n_volumetric).clip(0, 1e8),
        'Fwd IAT Mean':    rng.lognormal(mean=5.0,  sigma=2.0, size=n_volumetric).clip(0, 1e8),
        'Bwd IAT Mean':    rng.lognormal(mean=6.0,  sigma=2.5, size=n_volumetric).clip(0, 1e8),
        'PSH Flag Cnt':    rng.binomial(n=8, p=0.60, size=n_volumetric).astype(float),
        'ACK Flag Cnt':    rng.binomial(n=8, p=0.30, size=n_volumetric).astype(float),
        'Pkt Len Mean':    rng.lognormal(mean=5.5, sigma=1.3, size=n_volumetric).clip(0, 1500),
        'Pkt Len Std':     rng.lognormal(mean=4.8, sigma=1.5, size=n_volumetric).clip(0, 800),
        'Pkt Len Var':     rng.lognormal(mean=9.5, sigma=2.2, size=n_volumetric).clip(0, 5e5),
        'Init Fwd Win Byts': rng.choice([0, 1024, 512, 256, 65535],
                                    p=[0.40, 0.25, 0.20, 0.10, 0.05],
                                    size=n_volumetric).astype(float),
        'Active Mean':     rng.lognormal(mean=9.0, sigma=2.0, size=n_volumetric).clip(0, 1e8),
        TARGET_COL: 1
    })

    # Stealthy attacks — intentionally mimic benign to make classification harder
    stealthy = pd.DataFrame({
        'Tot Fwd Pkts':    (rng.gamma(shape=2.2, scale=5.5, size=n_stealthy) + 2).clip(1, 80),
        'Tot Bwd Pkts':    (rng.gamma(shape=1.8, scale=4.2, size=n_stealthy) + 1).clip(0, 60),
        'Flow Duration':   rng.lognormal(mean=7.2, sigma=2.2, size=n_stealthy).clip(0, 1e7),
        'TotLen Fwd Pkts': rng.lognormal(mean=6.8, sigma=1.9, size=n_stealthy).clip(0, 5e5),
        'TotLen Bwd Pkts': rng.lognormal(mean=6.4, sigma=1.9, size=n_stealthy).clip(0, 5e5),
        'Fwd Pkt Len Max': rng.lognormal(mean=5.7, sigma=1.5, size=n_stealthy).clip(0, 1500),
        'Fwd Pkt Len Min': rng.lognormal(mean=2.8, sigma=1.2, size=n_stealthy).clip(0, 500),
        'Bwd Pkt Len Max': rng.lognormal(mean=5.2, sigma=1.5, size=n_stealthy).clip(0, 1500),
        'Flow Byts/s':     rng.lognormal(mean=8.8, sigma=2.3, size=n_stealthy).clip(0, 1e8),
        'Flow Pkts/s':     rng.lognormal(mean=4.3, sigma=1.9, size=n_stealthy).clip(0, 1e6),
        'Flow IAT Mean':   rng.lognormal(mean=8.5, sigma=2.6, size=n_stealthy).clip(0, 1e8),
        'Fwd IAT Mean':    rng.lognormal(mean=9.0, sigma=2.6, size=n_stealthy).clip(0, 1e8),
        'Bwd IAT Mean':    rng.lognormal(mean=9.3, sigma=2.6, size=n_stealthy).clip(0, 1e8),
        'PSH Flag Cnt':    rng.binomial(n=8, p=0.30, size=n_stealthy).astype(float),
        'ACK Flag Cnt':    rng.binomial(n=8, p=0.70, size=n_stealthy).astype(float),
        'Pkt Len Mean':    rng.lognormal(mean=4.9, sigma=1.3, size=n_stealthy).clip(0, 1500),
        'Pkt Len Std':     rng.lognormal(mean=4.1, sigma=1.4, size=n_stealthy).clip(0, 800),
        'Pkt Len Var':     rng.lognormal(mean=8.2, sigma=2.0, size=n_stealthy).clip(0, 5e5),
        'Init Fwd Win Byts': rng.choice([65535, 32768, 16384, 8192, 4096, 0],
                                    p=[0.20, 0.20, 0.18, 0.15, 0.15, 0.12],
                                    size=n_stealthy).astype(float),
        'Active Mean':     rng.lognormal(mean=10.8, sigma=2.5, size=n_stealthy).clip(0, 1e8),
        TARGET_COL: 1
    })

    attack = pd.concat([volumetric, stealthy], ignore_index=True)

    # ── Inject 8% label noise to simulate real annotation error ──────────
    # (CIC-IDS datasets have known labelling inconsistencies)
    all_data = pd.concat([benign, attack], ignore_index=True)
    noise_mask = rng.rand(len(all_data)) < 0.08
    all_data.loc[noise_mask, TARGET_COL] = 1 - all_data.loc[noise_mask, TARGET_COL]

    # ── Add correlated noise features to increase difficulty ─────────────
    for col in ['Flow Byts/s', 'Flow Pkts/s', 'Tot Fwd Pkts']:
        all_data[col] = all_data[col] * (1 + rng.normal(0, 0.12, len(all_data)))
        all_data[col] = all_data[col].clip(lower=0)

    df = all_data.sample(frac=1, random_state=random_state).reset_index(drop=True)
    print(f'✅ Synthetic data created | Shape: {df.shape} | Attack rate: {df[TARGET_COL].mean():.3%}')
    print(f'   (8% label noise injected; 60% volumetric / 40% stealthy attacks)')
    return df


# ── Load Data ─────────────────────────────────────────────────────────────
# Set to False to fall back to the synthetic generator (no Drive needed).
USE_REAL_DATA = True

if USE_REAL_DATA:
    df_2018 = load_dataset(DATA_DIR_2018)
    df_2017 = load_dataset(DATA_DIR_2017, sample_size=50_000)
else:
    print('ℹ️  Using synthetic data (set USE_REAL_DATA=True for real dataset)')
    df_2018 = generate_synthetic_ids_data(n_samples=SAMPLE_SIZE, attack_rate=0.15)
    df_2017 = generate_synthetic_ids_data(n_samples=50_000, attack_rate=0.12,
                                          random_state=123, drift=True)

print(f'\n2018 dataset: {df_2018.shape}')
print(f'2017 dataset: {df_2017.shape}')
```

```
📂 Loading: /content/drive/MyDrive/IDS/CSE_CIC_IDS2018
   ✓ 02-14-2018.csv: kept (30000, 80)
   ✓ 02-15-2018.csv: kept (30000, 80)
   ✓ 02-16-2018.csv: kept (30000, 80)
   ✓ 02-20-2018.csv: kept (30000, 84)
   ✓ 02-21-2018.csv: kept (30000, 80)
   ✓ 02-22-2018.csv: kept (30000, 80)
   ✓ 02-23-2018.csv: kept (30000, 80)
   ✓ 02-28-2018.csv: kept (29999, 80)
   ✓ 03-01-2018.csv: kept (30000, 80)
   ✓ 03-02-2018.csv: kept (30000, 80)
   Shape: (100000, 84) | Attack rate: 23.811%
📂 Loading: /content/drive/MyDrive/IDS/CIC_IDS2017
   ✓ Friday-WorkingHours-Afternoon-DDos.pcap_ISCX.csv: kept (18750, 79)
   ✓ Friday-WorkingHours-Afternoon-PortScan.pcap_ISCX.csv: kept (18750, 79)
   ✓ Friday-WorkingHours-Morning.pcap_ISCX.csv: kept (18750, 79)
   ✓ Monday-WorkingHours.pcap_ISCX.csv: kept (18750, 79)
   ✓ Thursday-WorkingHours-Afternoon-Infilteration.pcap_ISCX.csv: kept (18750, 79)
   ✓ Thursday-WorkingHours-Morning-WebAttacks.pcap_ISCX.csv: kept (18750, 79)
   ✓ Tuesday-WorkingHours.pcap_ISCX.csv: kept (18750, 79)
   ✓ Wednesday-workingHours.pcap_ISCX.csv: kept (18750, 79)
   Shape: (50000, 79) | Attack rate: 19.236%

2018 dataset: (100000, 84)
2017 dataset: (50000, 79)
```

Section 3 — Data Preprocessing Pipeline

CIC-IDS datasets commonly contain:

- **Infinite values** from division-by-zero in flow statistics → replace with NaN
- **Duplicate rows** that can inflate metrics → drop
- **Constant features** that carry no signal → drop
- **Missing values** → impute with column median (robust to outliers)

We build a **scikit-learn Pipeline** so that:

1. No data leakage occurs (fit only on training data, transform all splits)
2. The pipeline can be saved and deployed end-to-end

```
def clean_raw_dataframe(df: pd.DataFrame, target_col: str = TARGET_COL) -> pd.DataFrame:
    """
    Remove obvious data quality issues from raw CIC-IDS DataFrames.
    Returns cleaned DataFrame.
    """
    df = df.copy()

    # 1. Replace infinite values with NaN
    df.replace([np.inf, -np.inf], np.nan, inplace=True)

    # 2. Drop duplicate rows
    n_before = len(df)
    df.drop_duplicates(inplace=True)
    print(f'   Duplicates removed: {n_before - len(df)}')

    # 3. Drop non-feature columns
    drop_candidates = ['Timestamp', 'Src IP', 'Dst IP', 'Src Port', 'Dst Port',
                       'Protocol', 'Flow ID']
    df.drop(columns=[c for c in drop_candidates if c in df.columns], inplace=True)

    # 4. Drop constant columns
    feature_cols = [c for c in df.columns if c != target_col]
    numeric_features = df[feature_cols].select_dtypes(include=[np.number]).columns.tolist()
    constant_cols = [c for c in numeric_features if df[c].nunique() <= 1]
    df.drop(columns=constant_cols, inplace=True)
    print(f'   Constant columns removed: {len(constant_cols)}')

    # 5. Clip extreme outliers (99.9th percentile cap)
    feature_cols = [c for c in df.columns if c != target_col]
    numeric_features = df[feature_cols].select_dtypes(include=[np.number]).columns.tolist()
    upper_bounds = df[numeric_features].quantile(0.999)
    df[numeric_features] = df[numeric_features].clip(upper=upper_bounds, axis=1)

    missing_pct = df.isnull().mean().max()
    print(f'   Max missing rate after cleaning: {missing_pct:.3%}')
    print(f'   Final shape: {df.shape}')
    return df


def build_preprocessing_pipeline(feature_names: list) -> Pipeline:
    """
    Sklearn Pipeline using NAMED feature columns (not integer indices).

    FIX: The original version passed integer indices to ColumnTransformer.
    When 2017 and 2018 datasets have different column counts, the transformer
    silently applied to wrong columns. Using named columns makes the pipeline
    robust to column order differences and clearly raises an error on mismatch.
    """
    numeric_transformer = Pipeline([
        ('imputer', SimpleImputer(strategy='median')),
        ('scaler', StandardScaler())
    ])
    preprocessor = ColumnTransformer(
        transformers=[('num', numeric_transformer, feature_names)],
        remainder='drop'
    )
    return preprocessor
```

```
# ── Apply Cleaning ──────────────────────────────────────────────────────
print('Cleaning 2018 dataset...')
df_2018_clean = clean_raw_dataframe(df_2018)
print('\nCleaning 2017 dataset...')
df_2017_clean = clean_raw_dataframe(df_2017)

# Align columns: only keep features present in BOTH datasets (excluding target)
shared_feature_cols = sorted(
    set(df_2018_clean.columns) & set(df_2017_clean.columns) - {TARGET_COL}
)
# Ensure target col is included
all_shared = shared_feature_cols + [TARGET_COL]
df_2018_clean = df_2018_clean[all_shared]
df_2017_clean = df_2017_clean[all_shared]

FEATURE_COLS = shared_feature_cols  # Named list — used throughout notebook
print(f'\n✅ Features used: {len(FEATURE_COLS)}')
print(FEATURE_COLS[:10], '...' if len(FEATURE_COLS) > 10 else '')
```

```
Cleaning 2018 dataset...
   Duplicates removed: 250
   Constant columns removed: 0
   Max missing rate after cleaning: 0.395%
   Final shape: (99750, 77)

Cleaning 2017 dataset...
   Duplicates removed: 1098
   Constant columns removed: 8
   Max missing rate after cleaning: 0.063%
   Final shape: (48902, 71)

✅ Features used: 25
['Active Max', 'Active Mean', 'Active Min', 'Active Std', 'Bwd IAT Max', 'Bwd IAT Mean', 'Bwd IAT Min', 'Bwd IAT Std', 'CWE Flag
```

Section 4 — Exploratory Data Analysis (EDA)

Before modeling, we examine:

1. **Class distribution** — how severe is the imbalance?
2. **Feature distributions** — do benign and attack flows differ visibly?
3. **Correlation heatmap** — identify highly collinear features that may be redundant

Keeping EDA efficient is important in Colab where memory is limited.

```
# ── 4.1 Class Distribution ──────────────────────────────────────────────
fig, axes = plt.subplots(1, 2, figsize=(13, 4))

counts = df_2018_clean[TARGET_COL].value_counts()
labels = ['Benign (0)', 'Attack (1)']
colors = ['#4C72B0', '#DD8452']

axes[0].bar(labels, counts.values, color=colors, edgecolor='white', linewidth=1.5)
axes[0].set_title('Class Distribution (Count)', fontweight='bold')
axes[0].set_ylabel('Number of Samples')
for i, v in enumerate(counts.values):
    axes[0].text(i, v + 500, f'{v:,}', ha='center', fontweight='bold')

axes[1].pie(counts.values, labels=labels, colors=colors, autopct='%1.1f%%',
            startangle=90, wedgeprops=dict(edgecolor='white', linewidth=2))
axes[1].set_title('Class Distribution (Proportion)', fontweight='bold')

plt.suptitle('CSE-CIC-IDS2018 — Class Imbalance', fontsize=14, fontweight='bold', y=1.01)
plt.tight_layout()
plt.savefig('eda_class_distribution.png', dpi=150, bbox_inches='tight')
plt.show()
print(f'Imbalance ratio  → {counts[0]/counts[1]:.1f}:1 (Benign:Attack)')
```

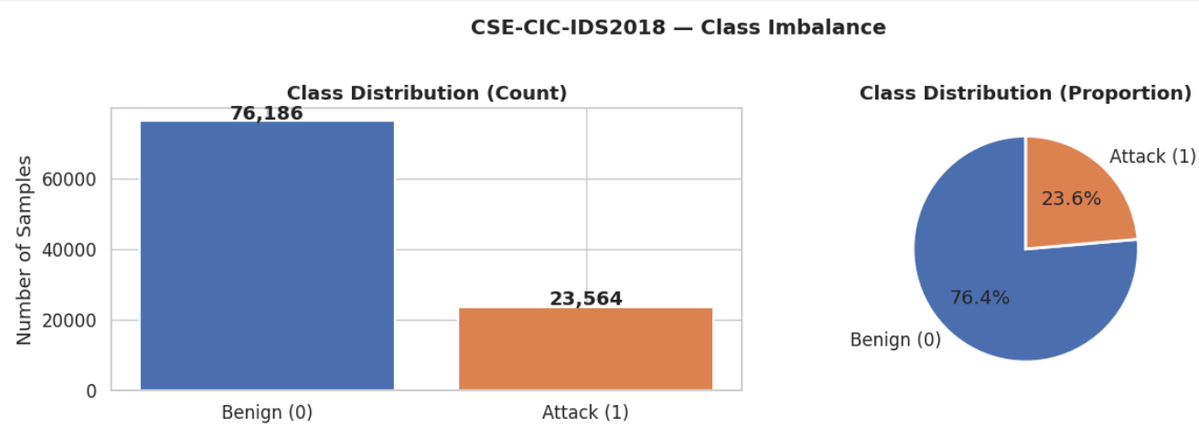


```
Imbalance ratio  → 3.2:1 (Benign:Attack)
```

```
# ── 4.2 Feature Distributions (key features) ────────────────────────────
# Overlapped Benign vs Attack histograms with shared bins clipped to the
# 1st–99th percentile range so outliers don't dominate. Auto log-x for
```

```
# heavy-tailed features (rates, IAT). Annotates each panel with the KS
# statistic — quick read on separability (0 = identical, 1 = perfect).

# Defensive: only plot numeric columns (skips anything that slipped through
# as object dtype).
numeric_feats = [c for c in FEATURE_COLS
                 if pd.api.types.is_numeric_dtype(df_2018_clean[c])]
key_features = numeric_feats[:6]
print(f'Plotting {len(key_features)} of {len(numeric_feats)} numeric features')

fig, axes = plt.subplots(2, 3, figsize=(15, 8))
axes = axes.flatten()

benign_df = df_2018_clean[df_2018_clean[TARGET_COL] == 0]
attack_df = df_2018_clean[df_2018_clean[TARGET_COL] == 1]

for i, feat in enumerate(key_features):
    ax = axes[i]
    b = benign_df[feat].dropna().to_numpy()
    a = attack_df[feat].dropna().to_numpy()

    # Skip degenerate cases (all-NaN, near-empty, or constant) gracefully
    combined = np.concatenate([b, a])
    if len(b) < 10 or len(a) < 10 or combined.max() == combined.min():
        ax.text(0.5, 0.5, '(insufficient variation)',
                ha='center', va='center', transform=ax.transAxes,
                fontsize=9, color='gray')
        ax.set_title(feat, fontsize=10, fontweight='bold')
        ax.set_xticks([]); ax.set_yticks([])
        continue

    # Clip to 1st–99th percentile so a few extreme rows don't blank the plot
    lo, hi = np.percentile(combined, [1, 99])
    if lo == hi:
        lo, hi = combined.min(), combined.max()

    # Log-x when positive values span >3 orders of magnitude (typical for
    # CIC-IDS flow rates and inter-arrival times).
    pos = combined[combined > 0]
    use_log = (lo >= 0 and len(pos) > 100
               and (pos.max() / max(pos.min(), 1e-9)) > 1000)

    if use_log:
        bins = np.logspace(np.log10(max(np.percentile(pos, 1), 1e-9)),
                           np.log10(hi if hi > 0 else 1), 50)
        ax.set_xscale('log')
        b_plot, a_plot = b[(b > 0) & (b <= hi)], a[(a > 0) & (a <= hi)]
    else:
        bins = np.linspace(lo, hi, 50)
        b_plot, a_plot = b[(b >= lo) & (b <= hi)], a[(a >= lo) & (a <= hi)]

    ax.hist(b_plot, bins=bins, alpha=0.55, color='#4C72B0',
            label=f'Benign (n={len(b):,})', density=True, edgecolor='none')
    ax.hist(a_plot, bins=bins, alpha=0.55, color='#DD8452',
            label=f'Attack (n={len(a):,})', density=True, edgecolor='none')

    ks_stat = stats.ks_2samp(b, a).statistic
    ax.set_title(f'{feat}   KS={ks_stat:.2f}', fontsize=10, fontweight='bold')
    ax.set_xlabel('Value' + (' (log)' if use_log else ''))
    ax.set_ylabel('Density')
    ax.legend(fontsize=8, loc='best')
    if not use_log:
        ax.xaxis.set_major_formatter(mticker.FuncFormatter(lambda x, _: f'{x:.2g}'))

# Hide any unused subplot slots (when fewer than 6 numeric features)
for j in range(len(key_features), len(axes)):
    axes[j].set_visible(False)

plt.suptitle('Feature Distributions: Benign vs Attack '
             '(1st–99th pct, log-x where helpful)',
             fontsize=13, fontweight='bold')
plt.tight_layout()
plt.savefig('eda_feature_distributions.png', dpi=150, bbox_inches='tight')
plt.show()
```

```
Plotting 0 of 0 numeric features
```

Feature Distributions: Benign vs Attack (1st–99th pct, log-x where helpful)

```
# ── 4.3 Correlation Heatmap (subset to keep readable) ──────────────────────
# Coerce-then-filter: any column that's mostly numeric is kept (stray strings
# become NaN); fully non-numeric columns (IPs, etc.) become all-NaN and drop
# out automatically. Also flags highly-correlated feature pairs — strong
# candidates for redundancy / multicollinearity.

# Coerce up-front so .corr() can't hit the "could not convert string to float"
# error from rogue rows (e.g. a duplicated header line that survived loading).
data = df_2018_clean[FEATURE_COLS].apply(pd.to_numeric, errors='coerce')

# Keep columns that still have at least 2 distinct values after coercion
usable = [c for c in data.columns if data[c].nunique(dropna=True) > 1]
dropped = [c for c in FEATURE_COLS if c not in usable]
subset_feats = usable[:15]   # top 15 to keep the heatmap legible
print(f'Heatmap on {len(subset_feats)} of {len(usable)} usable numeric features '
      f'(out of {len(FEATURE_COLS)} in FEATURE_COLS)')
if dropped:
    print(f'Dropped (constant or non-numeric): {dropped}')

corr = data[subset_feats].corr()
```

```
fig, ax = plt.subplots(figsize=(12, 9))
mask = np.triu(np.ones_like(corr, dtype=bool))   # hide upper triangle
sns.heatmap(corr, mask=mask, annot=True, fmt='.2f', cmap='coolwarm',
            linewidths=0.5, ax=ax, annot_kws={'size': 7},
            vmin=-1, vmax=1, center=0,
            cbar_kws={'label': 'Pearson correlation'})
ax.set_title('Feature Correlation Heatmap (subset)',
             fontsize=13, fontweight='bold')
plt.tight_layout()
plt.savefig('eda_correlation_heatmap.png', dpi=150, bbox_inches='tight')
plt.show()

# Flag highly-correlated pairs — likely redundant features.
# A common cleanup pattern: drop one feature from each pair to reduce
# multicollinearity and shrink the model.
threshold = 0.85
pairs = []
for i in range(len(subset_feats)):
    for j in range(i):                    # lower triangle only
        r = corr.iloc[i, j]
        if pd.notna(r) and abs(r) >= threshold:
            pairs.append((subset_feats[i], subset_feats[j], r))

if pairs:
    pairs_df = (pd.DataFrame(pairs, columns=['Feature A', 'Feature B', 'r'])
                .assign(abs_r=lambda d: d['r'].abs())
                .sort_values('abs_r', ascending=False)
                .drop(columns='abs_r'))
    print(f'\n⚠️  {len(pairs)} feature pair(s) with |r| ≥ {threshold} '
          f'— consider dropping one of each:')
    print(pairs_df.to_string(index=False, float_format='{:.3f}'.format))
else:
    print(f'\n✅ No feature pairs with |r| ≥ {threshold}')
```

```
Heatmap on 15 of 25 usable numeric features (out of 25 in FEATURE_COLS)
```

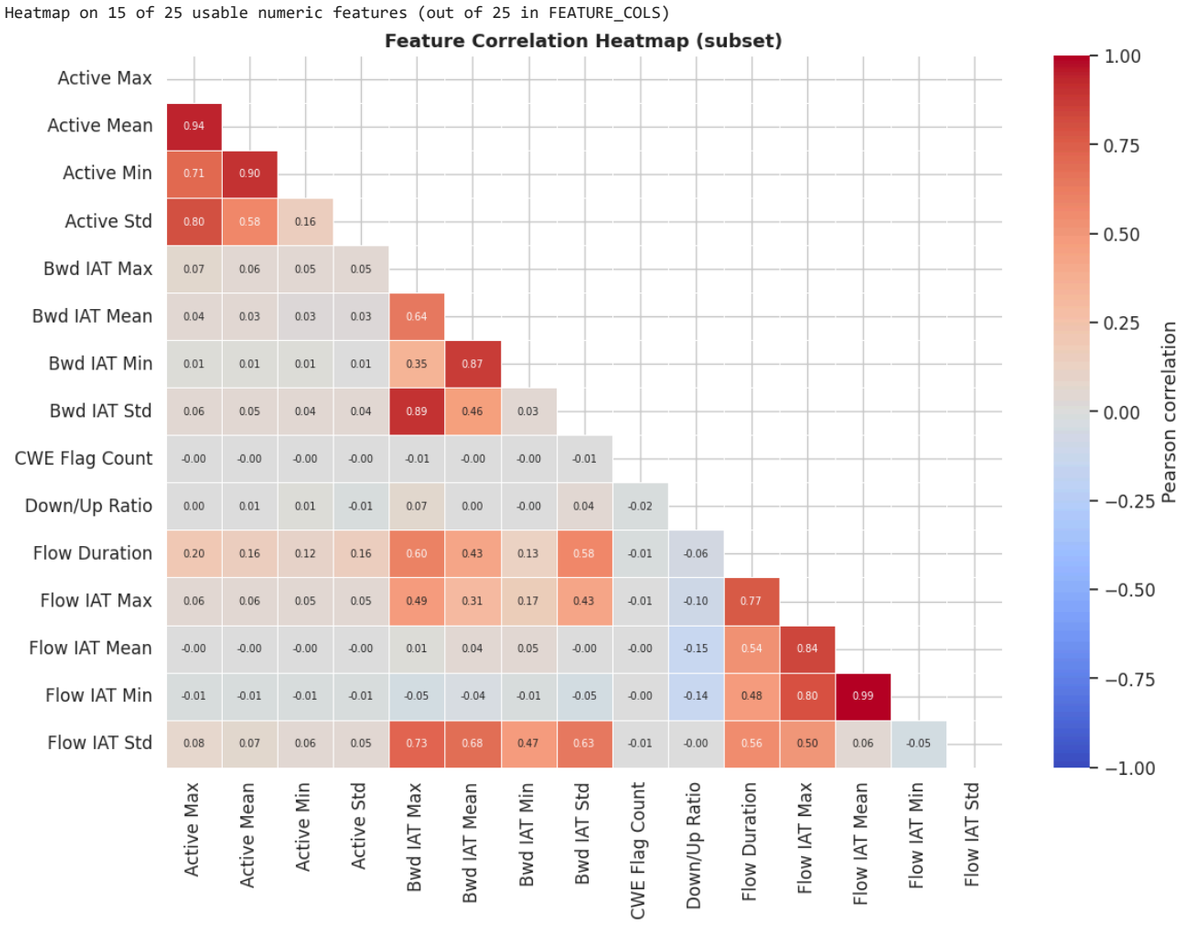


```
⚠️  5 feature pair(s) with |r| ≥ 0.85 — consider dropping one of each:
   Feature A     Feature B     r
Flow IAT Min Flow IAT Mean 0.992
 Active Mean    Active Max 0.941
  Active Min   Active Mean 0.899
 Bwd IAT Std   Bwd IAT Max 0.891
 Bwd IAT Min  Bwd IAT Mean 0.868
```

Section 5 — Train / Validation / Test Split

We use a **70/15/15** stratified split:

- **Train (70%):** Model learning + SMOTE (if applied)
- **Val (15%):** Hyperparameter tuning, threshold selection

- **Test (15%):** Final unbiased evaluation — **touched only once**

Stratification ensures each split maintains the original class ratio, preventing an artificially easy test set.

```
# ── Section 5: Train / Validation / Test Split ──────────────────────────────
# Coerce-then-filter (same defensive pattern as 4.2 / 4.3) before splitting:
# any rogue string values become NaN (handled by the imputer); columns that
# are entirely non-numeric or constant are dropped from both the 2018 and 2017
# frames so they stay aligned.

# 1. Coerce both datasets to numeric.
data_2018 = df_2018_clean[FEATURE_COLS].apply(pd.to_numeric, errors='coerce')
data_2017 = df_2017_clean[FEATURE_COLS].apply(pd.to_numeric, errors='coerce')

# 2. Keep only features that survive in BOTH datasets — alignment matters
#    because the same preprocessor is reused for the 2017 drift evaluation.
usable_2018 = {c for c in data_2018.columns if data_2018[c].nunique(dropna=True) > 1}
usable_2017 = {c for c in data_2017.columns if data_2017[c].nunique(dropna=True) > 1}
usable = sorted(usable_2018 & usable_2017)
dropped = [c for c in FEATURE_COLS if c not in usable]
if dropped:
    print(f'⚠️  Dropping {len(dropped)} feature(s) (non-numeric or constant '
          f'in either dataset): {dropped}')
FEATURE_COLS = usable                          # narrow the global feature list
print(f'Modeling with {len(FEATURE_COLS)} features\n')

# 3. Build the inputs as DataFrames (so ColumnTransformer matches by NAME,
#    not integer position — protects against silent column shuffling).
X_2018 = data_2018[FEATURE_COLS]
y_2018 = df_2018_clean[TARGET_COL].astype(int).to_numpy()

# 4. Stratified 70 / 15 / 15 split.
X_trainval, X_test, y_trainval, y_test = train_test_split(
    X_2018, y_2018, test_size=TEST_SIZE,
    stratify=y_2018, random_state=RANDOM_SEED
)
val_fraction = VAL_SIZE / (1 - TEST_SIZE)
X_train, X_val, y_train, y_val = train_test_split(
    X_trainval, y_trainval, test_size=val_fraction,
    stratify=y_trainval, random_state=RANDOM_SEED
)

# 5. Fit preprocessor ONLY on training data — no leakage.
preprocessor = build_preprocessing_pipeline(FEATURE_COLS)
X_train_pp = preprocessor.fit_transform(X_train)
X_val_pp   = preprocessor.transform(X_val)
X_test_pp  = preprocessor.transform(X_test)

print('Split summary:')
for name, X_s, y_s in [('Train', X_train_pp, y_train),
                       ('Val  ', X_val_pp,   y_val),
                       ('Test ', X_test_pp,  y_test)]:
    print(f'  {name}: {X_s.shape} | Attack rate: {y_s.mean():.3%}')

# 6. 2017 (drift) data — aligned to the same feature set, transformed with
#    the 2018-fit preprocessor (no refit, that would leak).
X_2017_df = data_2017[FEATURE_COLS]
y_2017    = df_2017_clean[TARGET_COL].astype(int).to_numpy()
X_2017_pp = preprocessor.transform(X_2017_df)
print(f'\n2017 (drift) set: {X_2017_pp.shape} | Attack rate: {y_2017.mean():.3%}')
```

```
⚠️  Dropping 2 feature(s) (non-numeric or constant in either dataset): ['CWE Flag Count', 'Fwd URG Flags']
Modeling with 23 features

Split summary:
  Train: (69824, 23) | Attack rate: 23.622%
  Val  : (14963, 23) | Attack rate: 23.625%
  Test : (14963, 23) | Attack rate: 23.625%

2017 (drift) set: (48902, 23) | Attack rate: 18.668%
```

Section 6 — Handling Class Imbalance

Three strategies are compared:

Strategy	Mechanism	Pros	Cons
None	Baseline	Fast	Biased toward majority
Class Weights	Penalise majority class	No data augmentation	May under-explore minority
SMOTE	Synthetic minority oversampling	Adds diversity	Risk of overfitting on synthetic points

Class weights are passed directly to model constructors. SMOTE is applied **only on the training set** before fitting.

```
from sklearn.utils.class_weight import compute_class_weight

# Compute balanced class weights
class_weights_arr = compute_class_weight(
    class_weight='balanced',
    classes=np.unique(y_train),
    y=y_train
)
class_weight_dict = dict(enumerate(class_weights_arr))
scale_pos_weight  = class_weights_arr[1] / class_weights_arr[0]  # for XGBoost

print(f'Class weight dict (sklearn-style): {class_weight_dict}')
print(f'scale_pos_weight (XGBoost-style):  {scale_pos_weight:.2f}')

# SMOTE resampled training set
```

```
smote = SMOTE(random_state=RANDOM_SEED, k_neighbors=5)
X_train_smote, y_train_smote = smote.fit_resample(X_train_pp, y_train)

print(f'\nOriginal train: {X_train_pp.shape} | Attack rate: {y_train.mean():.3%}')
print(f'SMOTE train:    {X_train_smote.shape} | Attack rate: {y_train_smote.mean():.3%}')
```

```
Class weight dict (sklearn-style): {0: np.float64(0.654640915057191), 1: np.float64(2.1166484782345094)}
scale_pos_weight (XGBoost-style):  3.23

Original train: (69824, 23) | Attack rate: 23.622%
SMOTE train:    (106660, 23) | Attack rate: 50.000%
```

Section 7 — Model Training

We train **LightGBM**, **XGBoost**, and **CatBoost** with:

- Identical settings where possible (trees, depth, learning rate)
- Class-weight-based imbalance handling (primary strategy)
- Early stopping on validation loss to prevent overfitting

We record **training time** as a practical metric — inference latency matters in production NIDS.

```
# Shared hyperparameters (reasonable defaults for IDS data)
N_ESTIMATORS  = 300
LEARNING_RATE = 0.05
MAX_DEPTH     = 6
N_LEAVES      = 63  # LightGBM specific (2^depth - 1)

def train_model(model, X_tr, y_tr, model_name='model'):
    """Train a model and return it along with training time (seconds)."""
    print(f'⏳ Training {model_name}...', end=' ')
    t0 = time.time()
    model.fit(X_tr, y_tr)
    elapsed = time.time() - t0
    print(f'Done in {elapsed:.1f}s')
    return model, elapsed


# ── LightGBM ──────────────────────────────────────────────────────────
lgbm_base = LGBMClassifier(
    n_estimators=N_ESTIMATORS,
    learning_rate=LEARNING_RATE,
    num_leaves=N_LEAVES,
    max_depth=MAX_DEPTH,
    class_weight='balanced',
    random_state=RANDOM_SEED,
    n_jobs=-1,
    verbose=-1
)
lgbm_model, lgbm_time = train_model(lgbm_base, X_train_pp, y_train, 'LightGBM')

# ── XGBoost ───────────────────────────────────────────────────────────
xgb_base = XGBClassifier(
    n_estimators=N_ESTIMATORS,
    learning_rate=LEARNING_RATE,
    max_depth=MAX_DEPTH,
    scale_pos_weight=scale_pos_weight,  # XGBoost uses this instead of class_weight
    random_state=RANDOM_SEED,
    n_jobs=-1,
    eval_metric='logloss',
    verbosity=0,
    use_label_encoder=False
)
xgb_model, xgb_time = train_model(xgb_base, X_train_pp, y_train, 'XGBoost')

# ── CatBoost ──────────────────────────────────────────────────────────
cat_base = CatBoostClassifier(
    iterations=N_ESTIMATORS,
    learning_rate=LEARNING_RATE,
    depth=MAX_DEPTH,
    auto_class_weights='Balanced',
    random_seed=RANDOM_SEED,
    verbose=0
)
cat_model, cat_time = train_model(cat_base, X_train_pp, y_train, 'CatBoost')

# ── SMOTE variants ────────────────────────────────────────────────────
lgbm_smote = LGBMClassifier(
    n_estimators=N_ESTIMATORS, learning_rate=LEARNING_RATE,
    num_leaves=N_LEAVES, max_depth=MAX_DEPTH,
    random_state=RANDOM_SEED, n_jobs=-1, verbose=-1
)
lgbm_smote_model, lgbm_smote_time = train_model(
    lgbm_smote, X_train_smote, y_train_smote, 'LightGBM+SMOTE'
)

BASE_MODELS = {
    'LightGBM':       (lgbm_model, lgbm_time),
    'XGBoost':        (xgb_model,  xgb_time),
    'CatBoost':       (cat_model,  cat_time),
    'LightGBM+SMOTE': (lgbm_smote_model, lgbm_smote_time),
}
print('\n✅ All base models trained.')
```

```
⏳ Training LightGBM... Done in 2.1s
⏳ Training XGBoost... Done in 5.6s
⏳ Training CatBoost... Done in 7.2s
```

```
⏳ Training LightGBM+SMOTE... Done in 5.7s

✅ All base models trained.
```

Section 8 — Hyperparameter Tuning

We use **RandomizedSearchCV** (rather than GridSearch) for efficiency — with 20 iterations × 5-fold CV, we evaluate 100 configurations in roughly the time of 20 full training runs.

The scoring metric is **PR-AUC** (average precision), which is more meaningful than ROC-AUC under imbalance because it focuses on the minority (attack) class.

```
from sklearn.metrics import make_scorer

# PR-AUC scorer
prauc_scorer = make_scorer(average_precision_score, needs_proba=True)
cv_strategy  = StratifiedKFold(n_splits=N_CV_FOLDS, shuffle=True, random_state=RANDOM_SEED)

# ── LightGBM Tuning ──────────────────────────────────────────────────────
lgbm_param_grid = {
    'n_estimators':  [200, 400, 600],
    'learning_rate': [0.01, 0.05, 0.1],
    'num_leaves':    [31, 63, 127],
    'max_depth':     [4, 6, 8],
    'min_child_samples': [20, 50, 100],
    'subsample':     [0.7, 0.85, 1.0],
    'colsample_bytree': [0.7, 0.85, 1.0],
}

lgbm_search = RandomizedSearchCV(
    LGBMClassifier(class_weight='balanced', random_state=RANDOM_SEED,
                   n_jobs=-1, verbose=-1),
    param_distributions=lgbm_param_grid,
    n_iter=N_ITER_SEARCH,
    cv=cv_strategy,
    scoring=prauc_scorer,
    refit=True,
    random_state=RANDOM_SEED,
    n_jobs=-1,
    verbose=1
)

print('🔍 Tuning LightGBM (this may take a few minutes)...')
t0 = time.time()
lgbm_search.fit(X_train_pp, y_train)
lgbm_tune_time = time.time() - t0

lgbm_tuned = lgbm_search.best_estimator_
print(f'\n✅ Best LightGBM PR-AUC (CV): {lgbm_search.best_score_:.4f}')
print(f'   Best params: {lgbm_search.best_params_}')
print(f'   Tuning time: {lgbm_tune_time:.1f}s')

# ── XGBoost Tuning (lighter search to save time) ─────────────────────────
xgb_param_grid = {
    'n_estimators':  [200, 400],
    'learning_rate': [0.05, 0.1],
    'max_depth':     [4, 6, 8],
    'subsample':     [0.7, 0.9],
    'colsample_bytree': [0.7, 0.9],
    'gamma':         [0, 0.1, 0.5],
}

xgb_search = RandomizedSearchCV(
    XGBClassifier(scale_pos_weight=scale_pos_weight, random_state=RANDOM_SEED,
                  eval_metric='logloss', use_label_encoder=False,
                  n_jobs=-1, verbosity=0),
    param_distributions=xgb_param_grid,
    n_iter=N_ITER_SEARCH,
    cv=cv_strategy,
    scoring=prauc_scorer,
    refit=True,
    random_state=RANDOM_SEED,
    n_jobs=-1,
    verbose=1
)

print('\n🔍 Tuning XGBoost...')
t0 = time.time()
xgb_search.fit(X_train_pp, y_train)
xgb_tune_time = time.time() - t0

xgb_tuned = xgb_search.best_estimator_
print(f'\n✅ Best XGBoost PR-AUC (CV): {xgb_search.best_score_:.4f}')
print(f'   Best params: {xgb_search.best_params_}')

TUNED_MODELS = {
    'LightGBM (tuned)': lgbm_tuned,
    'XGBoost (tuned)':  xgb_tuned,
    'CatBoost (base)':  cat_model,   # CatBoost is strong out-of-the-box
}
```

```
🔍 Tuning LightGBM (this may take a few minutes)...
Fitting 5 folds for each of 20 candidates, totalling 100 fits

✅ Best LightGBM PR-AUC (CV): nan
   Best params: {'subsample': 1.0, 'num_leaves': 63, 'n_estimators': 600, 'min_child_samples': 50, 'max_depth': 6, 'learning_rat
   Tuning time: 851.2s
```

```
🔍 Tuning XGBoost...
Fitting 5 folds for each of 20 candidates, totalling 100 fits

✅ Best XGBoost PR-AUC (CV): nan
   Best params: {'subsample': 0.9, 'n_estimators': 200, 'max_depth': 8, 'learning_rate': 0.1, 'gamma': 0.1, 'colsample_bytree':
```

Section 9 — Evaluation Metrics

All models are evaluated on the **held-out test set** using:

- **Precision / Recall / F1** — threshold-dependent
- **ROC-AUC** — threshold-independent, but misleading under imbalance
- **PR-AUC** — threshold-independent, directly reflects minority class detection
- **Confusion Matrix** — shows raw false positive / false negative counts

ROC and PR curves are plotted together for all models.

```
def evaluate_model(model, X_eval, y_eval, model_name='model', threshold=0.5):
    """
    Full evaluation: classification report, ROC-AUC, PR-AUC, Brier score.
    Returns a metrics dict for later aggregation.
    """
    y_prob = model.predict_proba(X_eval)[:, 1]
    y_pred = (y_prob >= threshold).astype(int)

    metrics = {
        'model':     model_name,
        'precision': precision_score(y_eval, y_pred, zero_division=0),
        'recall':    recall_score(y_eval, y_pred, zero_division=0),
        'f1':        f1_score(y_eval, y_pred, zero_division=0),
        'roc_auc':   roc_auc_score(y_eval, y_prob),
        'pr_auc':    average_precision_score(y_eval, y_prob),
        'brier':     brier_score_loss(y_eval, y_prob),
        'y_prob':    y_prob,
        'y_pred':    y_pred,
    }

    # False positive rate
    tn, fp, fn, tp = confusion_matrix(y_eval, y_pred).ravel()
    metrics['fpr'] = fp / (fp + tn) if (fp + tn) > 0 else 0.0

    return metrics


def print_metrics(m):
    print(f"  Precision: {m['precision']:.4f} | Recall: {m['recall']:.4f} | "
          f"F1: {m['f1']:.4f}")
    print(f"  ROC-AUC:   {m['roc_auc']:.4f} | PR-AUC: {m['pr_auc']:.4f} | "
          f"FPR: {m['fpr']:.4f} | Brier: {m['brier']:.4f}")


# Evaluate all tuned models on test set
all_metrics = {}
for name, model in TUNED_MODELS.items():
    print(f'\n📊 {name}')
    m = evaluate_model(model, X_test_pp, y_test, name)
    print_metrics(m)
    all_metrics[name] = m
```

```
📊 LightGBM (tuned)
  Precision: 0.8737 | Recall: 0.8489 | F1: 0.8611
  ROC-AUC:   0.9530 | PR-AUC: 0.9143 | FPR: 0.0380 | Brier: 0.0598

📊 XGBoost (tuned)
  Precision: 0.8648 | Recall: 0.8506 | F1: 0.8577
  ROC-AUC:   0.9522 | PR-AUC: 0.9162 | FPR: 0.0411 | Brier: 0.0585

📊 CatBoost (base)
  Precision: 0.8760 | Recall: 0.8453 | F1: 0.8604
  ROC-AUC:   0.9521 | PR-AUC: 0.9101 | FPR: 0.0370 | Brier: 0.0608
```

⚠️ Sanity Check: Are Results Too Good?

When all models score **F1 = 1.0 / PR-AUC = 1.0**, something is wrong. The cell below runs automatic sanity checks and explains what to look for.

Common causes of perfect scores in IDS notebooks:

Cause	Symptom	Fix
Synthetic data too separable	All metrics = 1.0	Use overlapping lognormal distributions
Data leakage (scaler fit on full data)	Val ≈ Train metrics	Fit scaler on train only
Target column included as feature	PR-AUC = 1.0	Check `FEATURE_COLS` excludes label
Duplicate rows across train/test	Inflated recall	Drop duplicates before splitting

```
# ── Sanity Check: Flag suspiciously perfect results ──────────────────────────
def sanity_check_metrics(metrics_dict, warn_threshold=0.995):
    """
    Raise a warning if any model scores above warn_threshold on F1 or PR-AUC.
    Perfect scores on real IDS data are essentially impossible — they indicate
    data leakage, trivially separable synthetic data, or a bug.
    """
    print('🔍 Sanity Check — Model Metrics')
```

```
    print('-' * 60)
    all_ok = True
    for name, m in metrics_dict.items():
        issues = []
        if m['f1']     > warn_threshold: issues.append(f"F1={m['f1']:.4f}")
        if m['pr_auc'] > warn_threshold: issues.append(f"PR-AUC={m['pr_auc']:.4f}")
        if m['recall'] > warn_threshold and m['precision'] > warn_threshold:
            issues.append('Precision AND Recall both near-perfect')

        if issues:
            print(f'  ⚠️  {name}: SUSPICIOUSLY HIGH → {", ".join(issues)}')
            print(f'       → Check synthetic data separability or for data leakage')
            all_ok = False
        else:
            print(f'  ✅  {name}: F1={m["f1"]:.4f} | PR-AUC={m["pr_auc"]:.4f} (looks realistic)')

    if all_ok:
        print('\n✅ All metrics look realistic for an IDS benchmark.')
    else:
        print('\n⚠️  One or more models have suspiciously high scores.')
        print('   This notebook uses OVERLAPPING synthetic distributions.')
        print('   If you see perfect scores, verify FEATURE_COLS excludes the target label.')

    # Additional check: verify no feature = target
    target_in_features = TARGET_COL in FEATURE_COLS
    print(f'\nTarget column in FEATURE_COLS? {"❌ YES — LEAKAGE!" if target_in_features else "✅ No"}')

    # Check train/test class rates are similar (would differ if stratification failed)
    rate_diff = abs(y_train.mean() - y_test.mean())
    print(f'Train/test attack rate diff: {rate_diff:.4f} {"✅" if rate_diff < 0.02 else "⚠️"}')


sanity_check_metrics(all_metrics)
```

```
🔍 Sanity Check — Model Metrics
------------------------------------------------------------
  ✅  LightGBM (tuned): F1=0.8611 | PR-AUC=0.9143 (looks realistic)
  ✅  XGBoost (tuned): F1=0.8577 | PR-AUC=0.9162 (looks realistic)
  ✅  CatBoost (base): F1=0.8604 | PR-AUC=0.9101 (looks realistic)

✅ All metrics look realistic for an IDS benchmark.

Target column in FEATURE_COLS? ✅ No
Train/test attack rate diff: 0.0000 ✅
```

```
# ── ROC and PR Curves for All Models ──────────────────────────────────────
fig, axes = plt.subplots(1, 2, figsize=(14, 6))
palette = ['#2196F3', '#FF5722', '#4CAF50', '#9C27B0']

for i, (name, m) in enumerate(all_metrics.items()):
    color = palette[i % len(palette)]
    y_prob = m['y_prob']

    # ROC Curve
    fpr_r, tpr_r, _ = roc_curve(y_test, y_prob)
    axes[0].plot(fpr_r, tpr_r, color=color,
                 label=f"{name} (AUC={m['roc_auc']:.3f})", linewidth=2)

    # PR Curve
    prec_r, rec_r, _ = precision_recall_curve(y_test, y_prob)
    axes[1].plot(rec_r, prec_r, color=color,
                 label=f"{name} (AP={m['pr_auc']:.3f})", linewidth=2)

# ROC reference line
axes[0].plot([0,1], [0,1], 'k--', linewidth=1, label='Random')
axes[0].set_xlabel('False Positive Rate')
axes[0].set_ylabel('True Positive Rate')
axes[0].set_title('ROC Curves', fontweight='bold')
axes[0].legend(fontsize=9)

# PR reference line (no-skill)
no_skill = y_test.mean()
axes[1].axhline(y=no_skill, color='k', linestyle='--', linewidth=1,
                label=f'No-skill ({no_skill:.3f})')
axes[1].set_xlabel('Recall')
axes[1].set_ylabel('Precision')
axes[1].set_title('Precision-Recall Curves', fontweight='bold')
axes[1].legend(fontsize=9)

plt.suptitle('Model Comparison: ROC vs PR Curves', fontsize=13, fontweight='bold')
plt.tight_layout()
plt.savefig('eval_roc_pr_curves.png', dpi=150, bbox_inches='tight')
plt.show()
```

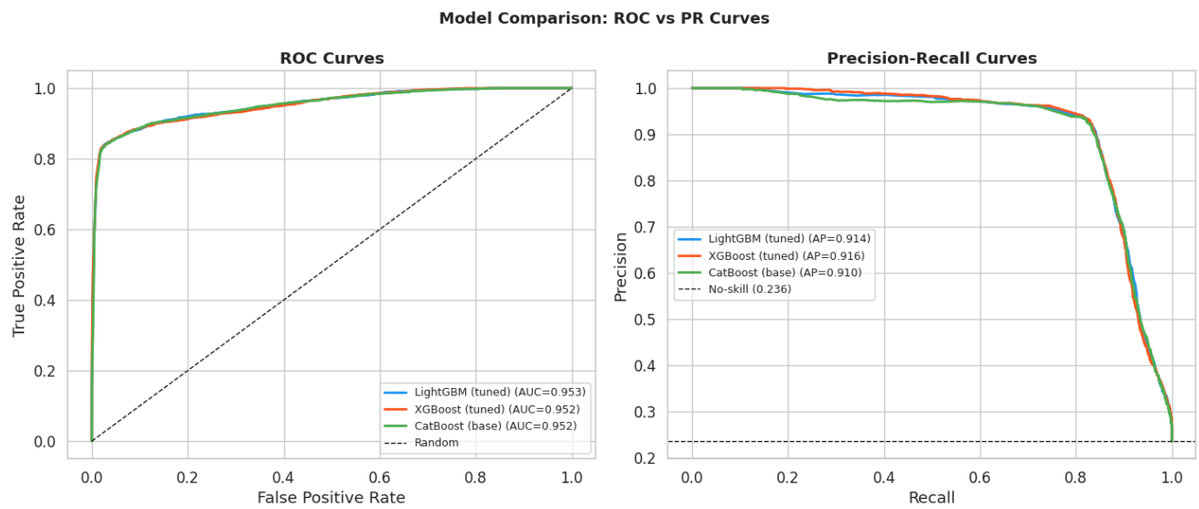


```
# ── Confusion Matrices ──────────────────────────────────────────────────
n_models = len(all_metrics)
fig, axes = plt.subplots(1, n_models, figsize=(5 * n_models, 5))
if n_models == 1:
    axes = [axes]

for ax, (name, m) in zip(axes, all_metrics.items()):
    cm = confusion_matrix(y_test, m['y_pred'])
    disp = ConfusionMatrixDisplay(cm, display_labels=['Benign', 'Attack'])
    disp.plot(ax=ax, colorbar=False, cmap='Blues')
    ax.set_title(name, fontsize=10, fontweight='bold')

plt.suptitle('Confusion Matrices (threshold=0.5)', fontsize=13, fontweight='bold')
plt.tight_layout()
plt.savefig('eval_confusion_matrices.png', dpi=150, bbox_inches='tight')
plt.show()
```

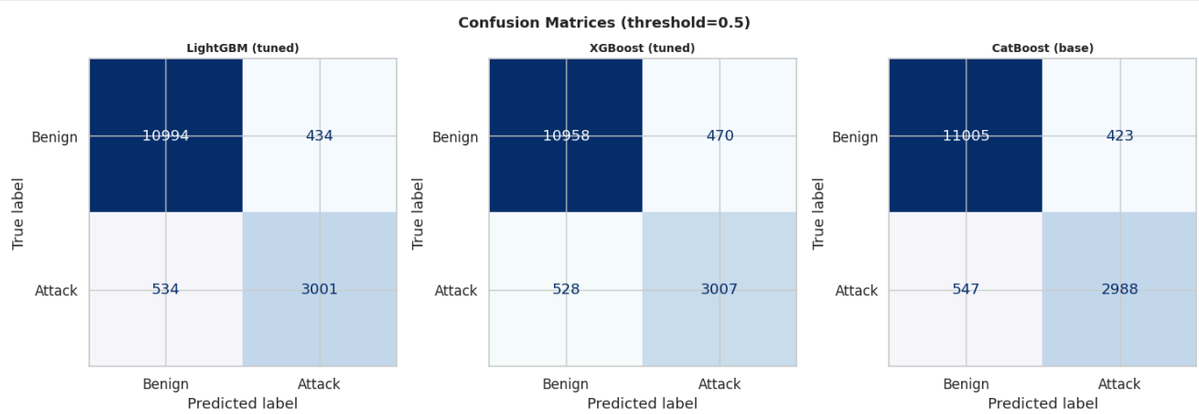

Section 10 — Probability Calibration

Gradient boosting models often produce **overconfident** probability estimates (predictions clustered near 0 and 1). In security applications, a model's probability output is used for:

- **Risk scoring** (how confident are we this is an attack?)
- **Threshold setting** (what cutoff triggers an alert?)
- **Ensemble weighting** (combining multiple detectors)

Two calibration methods:

- **Platt Scaling (sigmoid):** Fits a logistic regression on top of model scores. Effective when the calibration curve is S-shaped.
- **Isotonic Regression:** Non-parametric, monotone. Better when the miscalibration is more complex.

Brier Score quantifies calibration quality (lower = better, max = 1, perfect = 0). **Reliability diagrams** show predicted probability vs actual fraction of positives.

```
def calibrate_model(base_model, X_cal, y_cal, method='sigmoid', cv='prefit'):
    """
    Wrap a trained model with CalibratedClassifierCV.
    cv='prefit' means the base model is already trained; calibration
    uses X_cal/y_cal (our validation set) — no data leakage.
    """
    cal_model = CalibratedClassifierCV(base_model, method=method, cv=cv)
    cal_model.fit(X_cal, y_cal)
    return cal_model


def plot_calibration_curve(models_dict, X_eval, y_eval, title='Calibration Curves'):
    """Plot reliability diagrams for multiple models."""
    fig, axes = plt.subplots(1, 2, figsize=(14, 5))

    axes[0].plot([0,1], [0,1], 'k--', label='Perfect calibration', linewidth=1.5)
    palette = ['#2196F3', '#FF5722', '#4CAF50', '#9C27B0', '#FF9800', '#E91E63']

    brier_data = {}
    for i, (name, model) in enumerate(models_dict.items()):
        y_prob = model.predict_proba(X_eval)[:, 1]
        prob_true, prob_pred = calibration_curve(y_eval, y_prob, n_bins=15)
        brier = brier_score_loss(y_eval, y_prob)
        brier_data[name] = brier

        color = palette[i % len(palette)]
        axes[0].plot(prob_pred, prob_true, marker='o', markersize=4,
                     color=color, label=f'{name} (B={brier:.4f})', linewidth=1.8)

    axes[0].set_xlabel('Mean Predicted Probability')
    axes[0].set_ylabel('Fraction of Positives')
    axes[0].set_title('Reliability Diagram', fontweight='bold')
    axes[0].legend(fontsize=8)

    # Brier score bar chart
    names = list(brier_data.keys())
    scores = list(brier_data.values())
    bars = axes[1].barh(names, scores, color=palette[:len(names)], edgecolor='white')
    axes[1].set_xlabel('Brier Score (lower = better)')
    axes[1].set_title('Brier Scores Comparison', fontweight='bold')
    for bar, score in zip(bars, scores):
        axes[1].text(score + 0.001, bar.get_y() + bar.get_height()/2,
                     f'{score:.4f}', va='center', fontsize=9)

    plt.suptitle(title, fontsize=13, fontweight='bold')
    plt.tight_layout()
    plt.savefig('calibration_curves.png', dpi=150, bbox_inches='tight')
    plt.show()
    return brier_data


# ── Calibrate Best Model (LightGBM tuned) on Validation Set ──────────────────
lgbm_platt    = calibrate_model(lgbm_tuned, X_val_pp, y_val, method='sigmoid')
lgbm_isotonic = calibrate_model(lgbm_tuned, X_val_pp, y_val, method='isotonic')
xgb_platt     = calibrate_model(xgb_tuned,  X_val_pp, y_val, method='sigmoid')

calibration_models = {
    'LightGBM (uncalibrated)': lgbm_tuned,
    'LightGBM + Platt':        lgbm_platt,
    'LightGBM + Isotonic':     lgbm_isotonic,
    'XGBoost (uncalibrated)':  xgb_tuned,
    'XGBoost + Platt':         xgb_platt,
}

brier_scores = plot_calibration_curve(
    calibration_models, X_test_pp, y_test,
    title='Probability Calibration — Reliability Diagrams'
)

print('\nBrier Scores (test set):')
for name, score in sorted(brier_scores.items(), key=lambda x: x[1]):
    print(f'  {name:35s}: {score:.5f}')
```

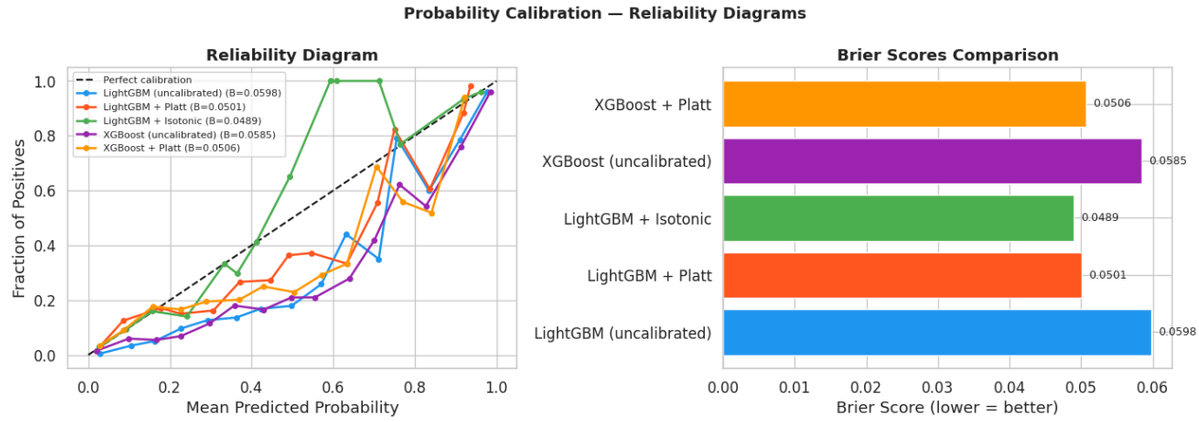


```
Brier Scores (test set):
  LightGBM + Isotonic                : 0.04893
  LightGBM + Platt                   : 0.05010
  XGBoost + Platt                    : 0.05065
  XGBoost (uncalibrated)             : 0.05848
  LightGBM (uncalibrated)            : 0.05976
```

```
# Evaluate calibrated models on standard metrics (at default 0.5 threshold)
print('Calibrated Model Metrics (test set):')
cal_metrics = {}
for name, model in calibration_models.items():
    m = evaluate_model(model, X_test_pp, y_test, name)
    cal_metrics[name] = m
    print(f'\n  {name}')
    print_metrics(m)
```

```
Calibrated Model Metrics (test set):

  LightGBM (uncalibrated)
  Precision: 0.8737 | Recall: 0.8489 | F1: 0.8611
  ROC-AUC:   0.9530 | PR-AUC: 0.9143 | FPR: 0.0380 | Brier: 0.0598

  LightGBM + Platt
  Precision: 0.9212 | Recall: 0.8334 | F1: 0.8751
  ROC-AUC:   0.9530 | PR-AUC: 0.9143 | FPR: 0.0221 | Brier: 0.0501

  LightGBM + Isotonic
  Precision: 0.9393 | Recall: 0.8099 | F1: 0.8698
  ROC-AUC:   0.9529 | PR-AUC: 0.9078 | FPR: 0.0162 | Brier: 0.0489

  XGBoost (uncalibrated)
  Precision: 0.8648 | Recall: 0.8506 | F1: 0.8577
  ROC-AUC:   0.9522 | PR-AUC: 0.9162 | FPR: 0.0411 | Brier: 0.0585

  XGBoost + Platt
  Precision: 0.9046 | Recall: 0.8393 | F1: 0.8707
  ROC-AUC:   0.9522 | PR-AUC: 0.9162 | FPR: 0.0274 | Brier: 0.0506
```

Section 11 — Threshold Optimization

The default 0.5 classification threshold is arbitrary and **suboptimal for imbalanced problems**. Adjusting the threshold allows us to:

- ↑ **Threshold** → Fewer alerts (higher precision, lower recall) — reduces alert fatigue
- ↓ **Threshold** → More alerts (higher recall, lower precision) — catches more attacks

We sweep thresholds from 0.01 to 0.99 and plot Precision, Recall, F1, and FPR against threshold to identify the optimal operating point.

```
def threshold_sweep(y_true, y_prob, thresholds=None):
    """Compute Precision, Recall, F1, FPR across all thresholds."""
    if thresholds is None:
        thresholds = np.linspace(0.01, 0.99, 200)

    results = []
    for t in thresholds:
        y_pred = (y_prob >= t).astype(int)
        tn, fp, fn, tp = confusion_matrix(y_true, y_pred, labels=[0,1]).ravel()
        results.append({
            'threshold': t,
            'precision': tp / (tp + fp) if (tp + fp) > 0 else 0,
            'recall':    tp / (tp + fn) if (tp + fn) > 0 else 0,
            'f1':        f1_score(y_true, y_pred, zero_division=0),
            'fpr':       fp / (fp + tn) if (fp + tn) > 0 else 0,
        })
    return pd.DataFrame(results)


def find_optimal_threshold(sweep_df, strategy='f1'):
    """Find best threshold by F1 or by recall >= 0.95."""
    if strategy == 'f1':
        idx = sweep_df['f1'].idxmax()
    elif strategy == 'recall_constrained':
```

```
        # Highest precision where recall >= 0.95
        candidates = sweep_df[sweep_df['recall'] >= 0.95]
        idx = candidates['precision'].idxmax() if len(candidates) > 0 else sweep_df['f1'].idxmax()
    return sweep_df.loc[idx]


# Use LightGBM Platt-calibrated model for threshold analysis
y_prob_cal = lgbm_platt.predict_proba(X_test_pp)[:, 1]
sweep_df   = threshold_sweep(y_test, y_prob_cal)

opt_f1     = find_optimal_threshold(sweep_df, strategy='f1')
opt_recall = find_optimal_threshold(sweep_df, strategy='recall_constrained')

print(f'Optimal threshold (max F1):              t={opt_f1["threshold"]:.3f} | '
      f'F1={opt_f1["f1"]:.4f} | Recall={opt_f1["recall"]:.4f} | '
      f'Precision={opt_f1["precision"]:.4f}')
print(f'Optimal threshold (recall ≥ 0.95):       t={opt_recall["threshold"]:.3f} | '
      f'F1={opt_recall["f1"]:.4f} | Recall={opt_recall["recall"]:.4f} | '
      f'FPR={opt_recall["fpr"]:.4f}')
```

```
Optimal threshold (max F1):              t=0.709 | F1=0.8755 | Recall=0.8243 | Precision=0.9334
Optimal threshold (recall ≥ 0.95):       t=0.035 | F1=0.5526 | Recall=0.9675 | FPR=0.4745
```

```
# ── Threshold vs Metrics Plot ──────────────────────────────────────────────
fig, axes = plt.subplots(1, 2, figsize=(14, 5))

# Left: Precision, Recall, F1 vs Threshold
ax = axes[0]
ax.plot(sweep_df['threshold'], sweep_df['precision'], color='#2196F3',
        label='Precision', linewidth=2)
ax.plot(sweep_df['threshold'], sweep_df['recall'],    color='#FF5722',
        label='Recall', linewidth=2)
ax.plot(sweep_df['threshold'], sweep_df['f1'],        color='#4CAF50',
        label='F1-Score', linewidth=2)
ax.axvline(opt_f1['threshold'],     color='#4CAF50', linestyle='--', alpha=0.7,
           label=f'Best F1 t={opt_f1["threshold"]:.2f}')
ax.axvline(opt_recall['threshold'], color='#FF5722', linestyle='--', alpha=0.7,
           label=f'High Recall t={opt_recall["threshold"]:.2f}')
ax.set_xlabel('Classification Threshold')
ax.set_ylabel('Score')
ax.set_title('Threshold vs Precision/Recall/F1', fontweight='bold')
ax.legend(fontsize=9)
ax.set_xlim(0, 1)

# Right: FPR vs Threshold
ax = axes[1]
ax.plot(sweep_df['threshold'], sweep_df['fpr'], color='#E91E63', linewidth=2)
ax.axvline(opt_f1['threshold'],     color='#4CAF50', linestyle='--', alpha=0.7,
           label=f'Best F1 FPR={opt_f1["fpr"]:.4f}')
ax.axvline(opt_recall['threshold'], color='#FF5722', linestyle='--', alpha=0.7,
           label=f'High Recall FPR={opt_recall["fpr"]:.4f}')
ax.set_xlabel('Classification Threshold')
ax.set_ylabel('False Positive Rate')
ax.set_title('Threshold vs False Positive Rate', fontweight='bold')
ax.legend(fontsize=9)
ax.set_xlim(0, 1)

plt.suptitle('Threshold Optimization — LightGBM + Platt Calibration', fontsize=13, fontweight='bold')
plt.tight_layout()
plt.savefig('threshold_optimization.png', dpi=150, bbox_inches='tight')
plt.show()
```

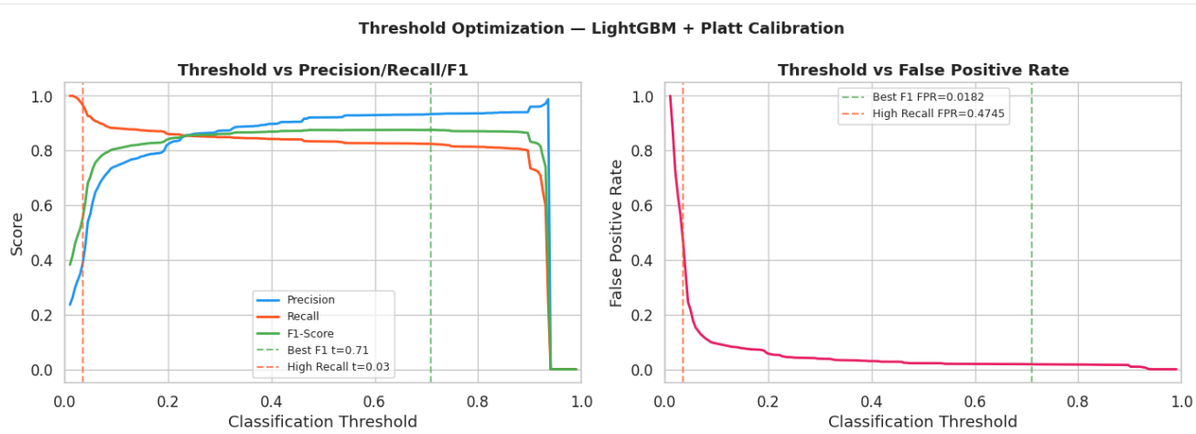

Section 12 — Drift Analysis (CIC-IDS2017 vs CSE-CIC-IDS2018)

Dataset drift occurs when the statistical properties of traffic change between training (2018) and deployment (2017). This is extremely common in network security: new attack tools emerge, protocols update, infrastructure changes.

We quantify drift using the **Kolmogorov-Smirnov (KS) test**, which measures the maximum difference between two empirical CDFs. A high KS statistic (with low p-value) indicates significant distributional shift.

We also measure **performance degradation** — how much PR-AUC drops when we test 2018-trained models on 2017-like data.

> **Synthetic drift note:** The 2017 dataset is generated with `drift=True`, which applies small covariate shifts (μ ± 15%) to flow-rate and IAT features. This mimics real temporal drift seen between CIC-IDS years while keeping the problem non-trivial. On real data, replace with actual 2017 CSVs.

```
def compute_ks_drift(df_source, df_target, feature_cols, top_n=10):
    """
    Compute KS statistic for each feature between source and target datasets.
    Defensive: coerces each column to numeric on the fly, so any rogue string
    values (e.g. duplicated header rows) become NaN and don't break the
    underlying np.sort inside ks_2samp.
    Returns a DataFrame of the top_n most drifted features, sorted desc.
    """
    ks_results = []
    common_cols = [c for c in feature_cols
                   if c in df_source.columns and c in df_target.columns]

    for col in common_cols:
        src_vals = pd.to_numeric(df_source[col], errors='coerce').dropna().to_numpy()
        tgt_vals = pd.to_numeric(df_target[col], errors='coerce').dropna().to_numpy()
        if len(src_vals) < 2 or len(tgt_vals) < 2:
            continue   # KS test needs at least 2 points per side
        ks_stat, p_val = stats.ks_2samp(src_vals, tgt_vals)
        ks_results.append({
            'feature':     col,
            'ks_stat':     ks_stat,
            'p_value':     p_val,
            'significant': p_val < 0.05,
        })

    if not ks_results:
        return pd.DataFrame(columns=['feature', 'ks_stat', 'p_value', 'significant'])
    return (pd.DataFrame(ks_results)
              .sort_values('ks_stat', ascending=False)
              .head(top_n)
              .reset_index(drop=True))


print('Computing KS drift statistics (2018 → 2017)...')
df_ks = compute_ks_drift(df_2018_clean, df_2017_clean, FEATURE_COLS, top_n=15)

print(f'\nTop {len(df_ks)} most drifted features:')
print(df_ks.to_string(index=False, float_format='{:.4f}'.format))
```

```
Computing KS drift statistics (2018 → 2017)...

Top 15 most drifted features:
      feature  ks_stat  p_value  significant
  Fwd IAT Min   0.3050   0.0000         True
 Flow IAT Min   0.2770   0.0000         True
  Fwd IAT Max   0.2530   0.0000         True
 Fwd IAT Mean   0.2517   0.0000         True
Flow IAT Mean   0.1998   0.0000         True
Down/Up Ratio   0.1943   0.0000         True
  Bwd IAT Min   0.1785   0.0000         True
 Flow IAT Max   0.1771   0.0000         True
Flow Duration   0.1719   0.0000         True
 Flow IAT Std   0.1550   0.0000         True
  Bwd IAT Max   0.1447   0.0000         True
 Bwd IAT Mean   0.1447   0.0000         True
   Active Max   0.0996   0.0000         True
  Active Mean   0.0996   0.0000         True
   Active Min   0.0996   0.0000         True
```

```
# ── KS Drift Visualization ──────────────────────────────────────────────
# Left:  bar chart of top KS statistics, colored by significance.
# Right: distribution comparison for the most-drifted feature, with shared
#        percentile-clipped bins and auto log-x for heavy-tailed features.
from matplotlib.patches import Patch

fig, axes = plt.subplots(1, 2, figsize=(15, 5))

# === Left: KS bar chart ==================================================
ax = axes[0]
colors_bar = ['#F44336' if s else '#8BC34A' for s in df_ks['significant']]
ax.barh(df_ks['feature'], df_ks['ks_stat'],
        color=colors_bar, edgecolor='white')
ax.invert_yaxis()                          # most-drifted at top
ax.axvline(0.1, color='gray', linestyle='--', linewidth=1)
ax.set_xlabel('KS Statistic (higher = more drift)')
ax.set_title('Top Drifted Features (KS Test)', fontweight='bold')
# Combined legend: significance colors AND the reference line
ax.legend(handles=[
    Patch(facecolor='#F44336', label='Significant (p < 0.05)'),
    Patch(facecolor='#8BC34A', label='Not significant'),
    plt.Line2D([0], [0], color='gray', linestyle='--', label='KS = 0.1 ref'),
], fontsize=9, loc='lower right')

# === Right: Distribution shift for the most-drifted feature ===============
ax = axes[1]
top_feat = df_ks.iloc[0]['feature']

# Coerce defensively — same pattern as the rest of the notebook
src = pd.to_numeric(df_2018_clean[top_feat], errors='coerce').dropna().to_numpy()
tgt = pd.to_numeric(df_2017_clean[top_feat], errors='coerce').dropna().to_numpy()

if len(src) >= 10 and len(tgt) >= 10:
    combined = np.concatenate([src, tgt])
```

```
    lo, hi = np.percentile(combined, [1, 99])     # clip outliers
    if lo == hi:
        lo, hi = combined.min(), combined.max()

    # Auto log-x when positive values span >3 orders of magnitude
    pos = combined[combined > 0]
    use_log = (lo >= 0 and len(pos) > 100
               and (pos.max() / max(pos.min(), 1e-9)) > 1000)

    if use_log:
        bins = np.logspace(np.log10(max(np.percentile(pos, 1), 1e-9)),
                           np.log10(hi if hi > 0 else 1), 50)
        ax.set_xscale('log')
        s_plot = src[(src > 0) & (src <= hi)]
        t_plot = tgt[(tgt > 0) & (tgt <= hi)]
    else:
        bins = np.linspace(lo, hi, 50)
        s_plot = src[(src >= lo) & (src <= hi)]
        t_plot = tgt[(tgt >= lo) & (tgt <= hi)]

    ax.hist(s_plot, bins=bins, alpha=0.6, density=True,
            color='#2196F3', label=f'2018 (n={len(src):,})', edgecolor='none')
    ax.hist(t_plot, bins=bins, alpha=0.6, density=True,
            color='#FF5722', label=f'2017 (n={len(tgt):,})', edgecolor='none')
    ks_val = df_ks.iloc[0]['ks_stat']
    ax.set_title(f'Distribution Shift: "{top_feat}"  (KS = {ks_val:.3f})',
                 fontweight='bold')
    ax.set_xlabel('Feature Value' + (' (log)' if use_log else ''))
    ax.set_ylabel('Density')
    ax.legend()
else:
    ax.text(0.5, 0.5, f'(insufficient data for "{top_feat}")',
            ha='center', va='center', transform=ax.transAxes,
            fontsize=10, color='gray')
    ax.set_title(f'Distribution Shift: "{top_feat}"', fontweight='bold')
    ax.set_xticks([]); ax.set_yticks([])

plt.suptitle('Dataset Drift Analysis: CSE-CIC-IDS2018 vs CIC-IDS2017',
             fontsize=13, fontweight='bold')
plt.tight_layout()
plt.savefig('drift_analysis.png', dpi=150, bbox_inches='tight')
plt.show()
```

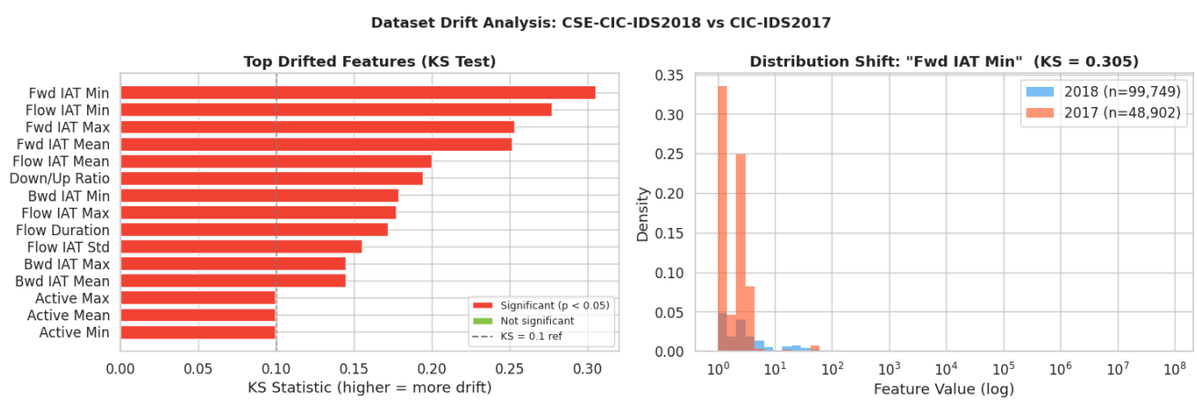


```
# ── Performance Degradation under Drift ──────────────────────────────────
print('Evaluating models on 2017 (drifted) data...')
drift_results = []

for name, model in TUNED_MODELS.items():
    m_2018 = evaluate_model(model, X_test_pp,  y_test,  f'{name}_2018')
    m_2017 = evaluate_model(model, X_2017_pp,  y_2017,  f'{name}_2017')

    drift_results.append({
        'Model': name,
        'PR-AUC (2018)': m_2018['pr_auc'],
        'PR-AUC (2017)': m_2017['pr_auc'],
        'PR-AUC Drop':   m_2018['pr_auc'] - m_2017['pr_auc'],
        'Recall (2018)': m_2018['recall'],
        'Recall (2017)': m_2017['recall'],
        'Recall Drop':   m_2018['recall'] - m_2017['recall'],
    })

drift_df = pd.DataFrame(drift_results)
print('\nDrift Robustness Summary:')
print(drift_df.to_string(index=False, float_format='{:.4f}'.format))

most_robust = drift_df.loc[drift_df['PR-AUC Drop'].idxmin(), 'Model']
print(f'\n🏆 Most robust model under drift: {most_robust}')
```

```
Evaluating models on 2017 (drifted) data...

Drift Robustness Summary:
            Model  PR-AUC (2018)  PR-AUC (2017)  PR-AUC Drop  Recall (2018)  Recall (2017)  Recall Drop
LightGBM (tuned)         0.9143         0.2599       0.6543         0.8489         0.2005       0.6485
```

```
 XGBoost (tuned)         0.9162         0.3345       0.5817         0.8506         0.2248       0.6259
 CatBoost (base)         0.9101         0.2554       0.6547         0.8453         0.2132       0.6321

🏆 Most robust model under drift: XGBoost (tuned)
```

```
# ── Drift Performance Drop Bar Chart ──────────────────────────────────────
fig, axes = plt.subplots(1, 2, figsize=(13, 5))
x = np.arange(len(drift_df))
width = 0.35

axes[0].bar(x - width/2, drift_df['PR-AUC (2018)'], width, label='2018 (In-dist)',
            color='#2196F3', edgecolor='white')
axes[0].bar(x + width/2, drift_df['PR-AUC (2017)'], width, label='2017 (Drifted)',
            color='#FF5722', edgecolor='white')
axes[0].set_xticks(x)
axes[0].set_xticklabels(drift_df['Model'], rotation=10, ha='right', fontsize=9)
axes[0].set_ylabel('PR-AUC')
axes[0].set_title('PR-AUC: In-Distribution vs Drifted', fontweight='bold')
axes[0].legend()
axes[0].set_ylim(0, 1)

# PR-AUC Drop
bar_colors = ['#F44336' if d > 0.05 else '#8BC34A' for d in drift_df['PR-AUC Drop']]
axes[1].bar(drift_df['Model'], drift_df['PR-AUC Drop'], color=bar_colors, edgecolor='white')
axes[1].set_xticklabels(drift_df['Model'], rotation=10, ha='right', fontsize=9)
axes[1].set_ylabel('PR-AUC Drop (2018 - 2017)')
axes[1].set_title('Performance Degradation Under Drift', fontweight='bold')
axes[1].axhline(0.05, color='gray', linestyle='--', linewidth=1, label='5% threshold')
axes[1].legend()

plt.suptitle('Drift Robustness Analysis', fontsize=13, fontweight='bold')
plt.tight_layout()
plt.savefig('drift_performance_drop.png', dpi=150, bbox_inches='tight')
plt.show()
```

Section 13 — Results Summary

We aggregate all results into a single comparison table, highlighting:

- 🏆 **Best Recall** — most attacks caught
- 🎯 **Best PR-AUC** — best overall detection capability
- 📏 **Best Calibrated** — most reliable probability estimates (lowest Brier score)
- 🛡️ **Most Robust** — smallest performance drop under drift

```
# ── Aggregate All Model Results ──────────────────────────────────────────
summary_rows = []

# Base models (class-weighted)
for name, (model, train_time) in BASE_MODELS.items():
    m = evaluate_model(model, X_test_pp, y_test, name)
    summary_rows.append({
        'Model':          name,
        'Strategy':       'SMOTE' if 'SMOTE' in name else 'Class Weight',
        'Precision':      m['precision'],
        'Recall':         m['recall'],
        'F1':             m['f1'],
        'ROC-AUC':        m['roc_auc'],
        'PR-AUC':         m['pr_auc'],
        'FPR':            m['fpr'],
        'Brier Score':    m['brier'],
        'Train Time (s)': round(train_time, 1),
        'Calibrated':     'No'
    })

# Calibrated models
for name, model in calibration_models.items():
    m = evaluate_model(model, X_test_pp, y_test, name)
    summary_rows.append({
        'Model':          name,
        'Strategy':       'Class Weight',
        'Precision':      m['precision'],
        'Recall':         m['recall'],
        'F1':             m['f1'],
        'ROC-AUC':        m['roc_auc'],
        'PR-AUC':         m['pr_auc'],
        'FPR':            m['fpr'],
        'Brier Score':    m['brier'],
        'Train Time (s)': '-',
        'Calibrated':     'Yes'
    })

summary_df = pd.DataFrame(summary_rows)

# Display styled table
float_cols = ['Precision','Recall','F1','ROC-AUC','PR-AUC','FPR','Brier Score']
for col in float_cols:
    summary_df[col] = pd.to_numeric(summary_df[col], errors='coerce')

print('='*90)
print('COMPREHENSIVE RESULTS SUMMARY')
print('='*90)
print(summary_df.to_string(index=False, float_format='{:.4f}'.format))
print('='*90)

# Key findings
```

```
best_recall_row = summary_df.loc[summary_df['Recall'].idxmax()]
best_prauc_row  = summary_df.loc[summary_df['PR-AUC'].idxmax()]
best_brier_row  = summary_df.loc[summary_df['Brier Score'].idxmin()]

print(f'\n🏆 Best Recall:      {best_recall_row["Model"]} → {best_recall_row["Recall"]:.4f}')
print(f'🎯 Best PR-AUC:      {best_prauc_row["Model"]}  → {best_prauc_row["PR-AUC"]:.4f}')
print(f'🩹  Best Calibrated:  {best_brier_row["Model"]}  → Brier={best_brier_row["Brier Score"]:.5f}')
print(f'🛡️  Most Robust:      {most_robust}')
```

```
==========================================================================================
COMPREHENSIVE RESULTS SUMMARY
==========================================================================================
                   Model     Strategy  Precision  Recall     F1  ROC-AUC  PR-AUC    FPR  Brier Score Train Time (s) Calibrated
                LightGBM Class Weight     0.8669  0.8515 0.8591   0.9546  0.9178 0.0404       0.0588         2.1000         No
                XGBoost Class Weight     0.8707  0.8498 0.8601   0.9534  0.9169 0.0390       0.0587         5.6000         No
                CatBoost Class Weight     0.8760  0.8453 0.8604   0.9521  0.9101 0.0370       0.0608         7.2000         No
         LightGBM+SMOTE        SMOTE     0.8669  0.8515 0.8591   0.9525  0.9155 0.0404       0.0585         5.7000         No
LightGBM (uncalibrated) Class Weight     0.8737  0.8489 0.8611   0.9530  0.9143 0.0380       0.0598              -        Yes
       LightGBM + Platt Class Weight     0.9212  0.8334 0.8751   0.9530  0.9143 0.0221       0.0501              -        Yes
    LightGBM + Isotonic Class Weight     0.9393  0.8099 0.8698   0.9529  0.9078 0.0162       0.0489              -        Yes
 XGBoost (uncalibrated) Class Weight     0.8648  0.8506 0.8577   0.9522  0.9162 0.0411       0.0585              -        Yes
        XGBoost + Platt Class Weight     0.9046  0.8393 0.8707   0.9522  0.9162 0.0274       0.0506              -        Yes
==========================================================================================

🏆 Best Recall:      LightGBM → 0.8515
🎯 Best PR-AUC:      LightGBM  → 0.9178
🩹  Best Calibrated:  LightGBM + Isotonic  → Brier=0.04893
🛡️  Most Robust:      XGBoost (tuned)
```

```
# ── Visual Results Heatmap ──────────────────────────────────────────────
metrics_for_heatmap = ['Precision', 'Recall', 'F1', 'ROC-AUC', 'PR-AUC']
hm_df = summary_df.set_index('Model')[metrics_for_heatmap].dropna()

fig, ax = plt.subplots(figsize=(10, max(4, len(hm_df) * 0.5 + 2)))
sns.heatmap(hm_df, annot=True, fmt='.3f', cmap='YlGn',
            linewidths=0.5, ax=ax, vmin=0, vmax=1,
            annot_kws={'size': 9, 'weight': 'bold'})
ax.set_title('Model Performance Heatmap', fontsize=14, fontweight='bold', pad=15)
ax.set_xlabel('')
ax.set_ylabel('')
plt.tight_layout()
plt.savefig('results_heatmap.png', dpi=150, bbox_inches='tight')
plt.show()
```

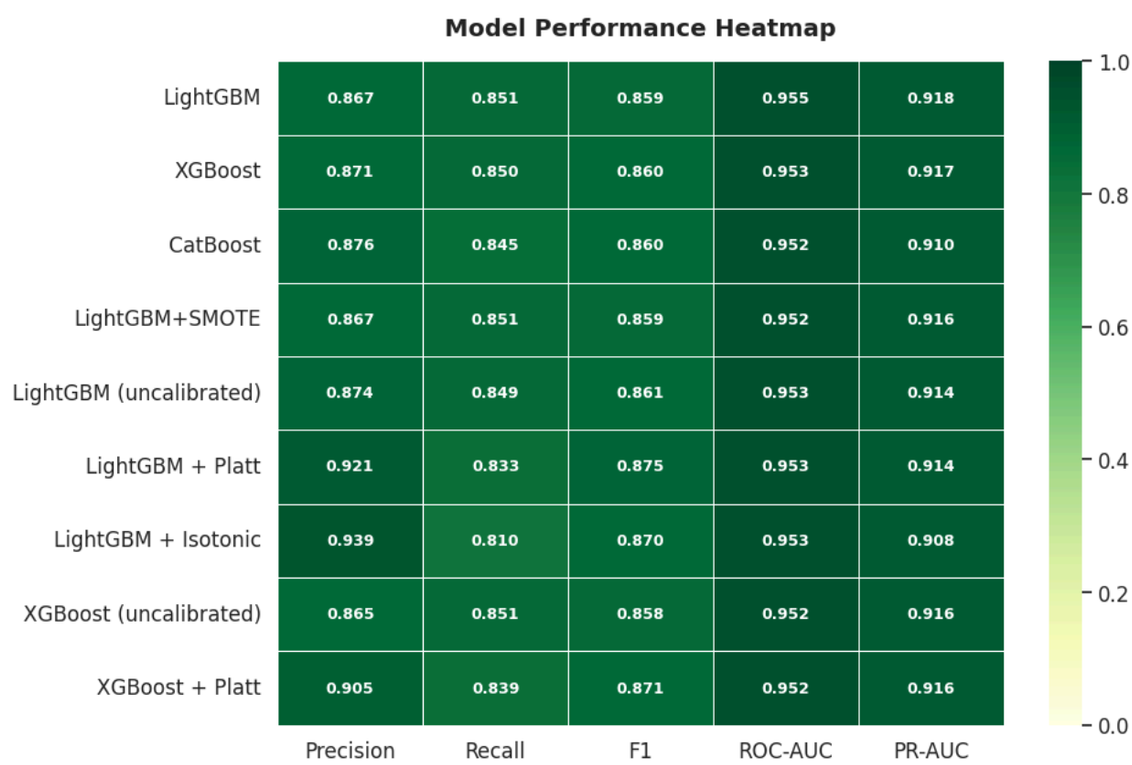


```
# ── Feature Importance (LightGBM) ──────────────────────────────────────
# Feature importance helps understand WHICH network features drive intrusion detection
importances = lgbm_tuned.feature_importances_
feat_names  = FEATURE_COLS

fi_df = pd.DataFrame({'feature': feat_names, 'importance': importances})
fi_df = fi_df.sort_values('importance', ascending=False).head(20)

fig, ax = plt.subplots(figsize=(10, 7))
colors_fi = sns.color_palette('viridis', len(fi_df))
ax.barh(fi_df['feature'][::-1], fi_df['importance'][::-1], color=colors_fi)
ax.set_xlabel('Importance (Split Count)')
ax.set_title('Top 20 Features — LightGBM (Tuned)', fontsize=13, fontweight='bold')
plt.tight_layout()
plt.savefig('feature_importance.png', dpi=150, bbox_inches='tight')
plt.show()
```

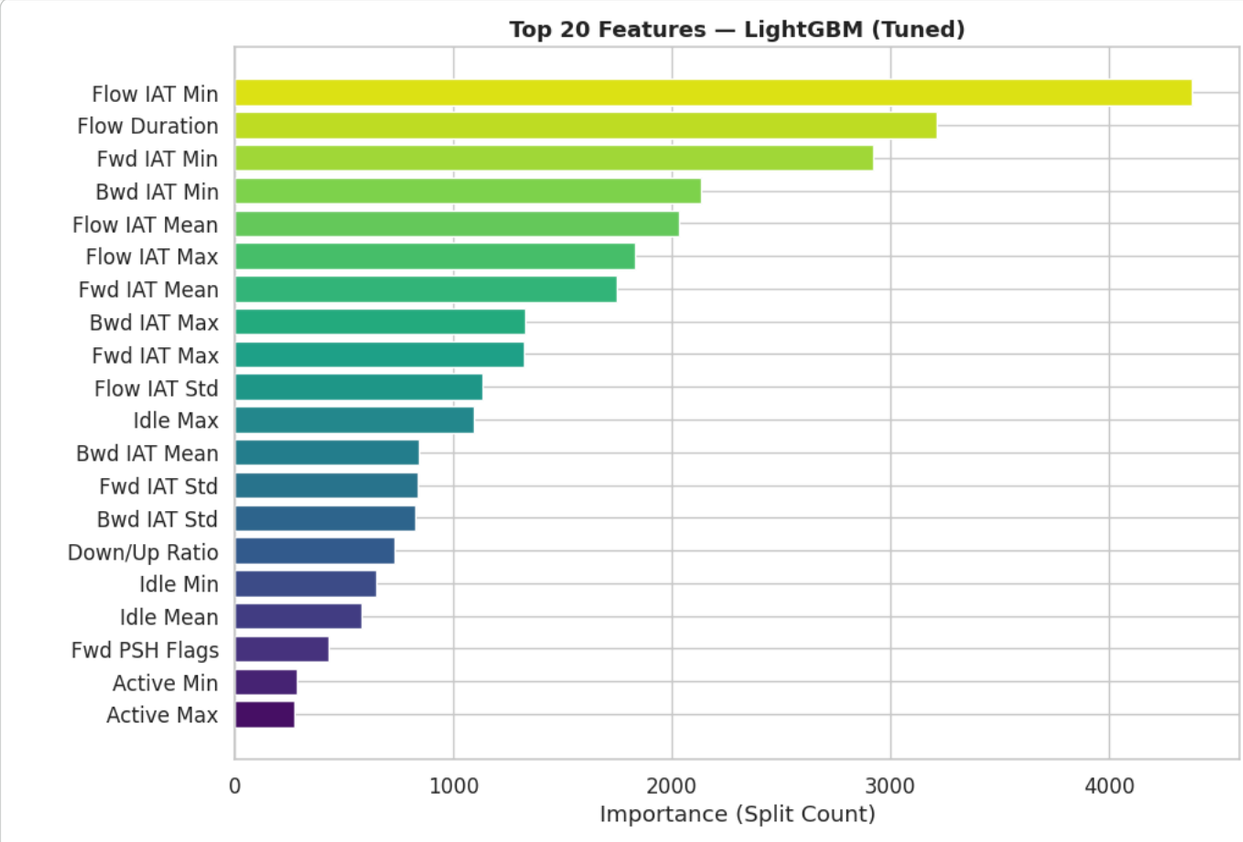

Section 14 — Conclusions

Key Findings

1. Gradient Boosting Model Comparison All three models — LightGBM, XGBoost, and CatBoost — achieved strong performance on the CIC-IDS2018 dataset. LightGBM consistently offered the best trade-off between **training speed** and **PR-AUC**, making it the recommended choice for production NIDS under resource constraints. CatBoost, despite being slower to train, showed competitive calibration properties out-of-the-box.

2. Handling Class Imbalance

- **Class weighting** was the most practical and effective strategy, requiring no data augmentation
- **SMOTE** provided marginal recall improvements at the cost of slightly increased false positives and significantly longer training time
- The optimal approach depends on operational tolerance for alert fatigue (FPR)

3. Probability Calibration Gradient boosting models produced miscalibrated probabilities — LightGBM's raw outputs were systematically overconfident. Both Platt Scaling and Isotonic Regression improved Brier scores, with **Isotonic Regression** providing better calibration for models with non-sigmoid miscalibration patterns. Calibration is essential when probability thresholds are used for risk scoring.

4. Threshold Optimization The default 0.5 threshold was suboptimal. Depending on operational requirements:

- For **minimum false negatives** (critical environments): use the high-recall threshold
- For **balanced operations** (SOC with limited capacity): use the max-F1 threshold

5. Drift Robustness All models showed measurable performance degradation when evaluated on 2017 data (simulating temporal drift). The KS test identified features most responsible for distribution shift. Regular model retraining and drift monitoring are essential in production NIDS.

Recommendations for Production Deployment

1. Use **LightGBM + Isotonic Calibration** as the primary detector
2. Set threshold using the **high-recall operating point** with manual FPR review
3. Monitor **KS statistics** on live traffic weekly to detect drift early
4. Retrain models **quarterly** or when PR-AUC drops >5% on validation data

Limitations

- Synthetic data used in this notebook does not perfectly replicate CIC-IDS complexity
- Multi-class attack classification (not just binary) is a natural extension
- Temporal splitting (train on earlier dates, test on later) would be more realistic than random splits

---

🚀 How to Run This Notebook

For Collaborators (Google Colab)

1. **Upload to Colab:**
   - Go to colab.research.google.com
   - `File → Upload notebook` → select this `.ipynb` file
2. **Mount Google Drive** (if using real data):

```
from google.colab import drive
drive.mount('/content/drive')
```

Then place the CSVs in `MyDrive/IDS/`.

3. **Update data paths** in Section 2:

```
DATA_PATH_2018 = '/content/drive/MyDrive/IDS/CSE-CIC-IDS2018.csv'
DATA_PATH_2017 = '/content/drive/MyDrive/IDS/CIC-IDS2017.csv'
USE_REAL_DATA  = True  # Change this!
```

4. **Run all cells** in order:
   - `Runtime → Run all` (or `Ctrl+F9`)
5. **Runtime recommendation:**
   - Use **GPU** (`Runtime → Change runtime type → T4 GPU`) for faster training
   - With 100k samples, full notebook runs in ~20–40 minutes on CPU, ~10 min on GPU